

Induced Metric Optimisers

March 26, 2026

Contents

1	Introduction	3
2	Setup and Notation	3
3	Fixed Diagonal Metric	4
3.1	Geometric Derivation	4
3.2	Practical Algorithm	5
3.3	Log-Loss Variant	7
3.4	Implementation	7
4	Learnable Scalar Metric	7
4.1	Construction	8
4.2	Metric Learning	8
4.3	Algorithm	10
4.4	Hyperparameters	11
4.5	Implementation	11
5	Learnable Diagonal Metric	11
5.1	Construction	11
5.2	Metric Learning	12
5.3	Mean-Centering	12
5.4	Algorithm	13
5.5	Implementation	14
6	Off-Diagonal Metric	14
6.1	Derivation	14
6.2	Algorithm	16
6.3	Implementation	16
7	Shared Implementation Details	17
7.1	Update Direction vs. Metric Direction	17
7.2	Computational Cost	18
7.3	Shared Hyperparameters	18
A	Repository Structure	18
A.1	Directory Layout	18
A.2	Running Sweeps	19
A.3	Available Optimisers	19
A.4	Analysis Notebooks	19

8	Geodesic Convexity Analysis	20
8.1	Background: Geodesic Convexity	20
8.2	Christoffel Symbols	20
8.3	Riemannian Hessian	21
8.4	Connection to the Gauss Equation	22
8.5	Off-Diagonal Metric	24
8.6	Learnable Metrics: Genuine Curvature Correction	24
8.7	Implications for the Optimiser	26

1 Introduction

Thomas [1] introduced a class of optimisers that precondition gradients using the Riemannian metric naturally induced when the loss landscape is embedded in higher-dimensional space. The key idea is to embed the N -dimensional parameter space into $\mathbb{R}^{N+1}$ via the loss function, equip the ambient space with a metric h , and pull h back to obtain a geometry-aware preconditioner. The resulting update automatically damps the effective learning rate in high-curvature regions—acting as a smooth form of gradient clipping—while adding only $O(N)$ overhead.

This document extends that framework in three directions (summarised in Table 1):

- (i) **Learnable scalar metric** (Section 4): the inverse base metric γ^{-1} is parameterised by a single learnable scalar s , allowing the overall metric scale to adapt during training.
- (ii) **Learnable diagonal metric** (Section 5): each parameter receives its own learnable scale s_i , enabling per-parameter metric adaptation with a principled gauge-fixing (mean-centering) to remove the degeneracy with the global scaling hyperparameter ξ .
- (iii) **Off-diagonal ambient metric** (Section 6): the ambient metric h in [1] has no coupling between the parameter dimensions and the loss dimension (the off-diagonal blocks are zero). We relax this by introducing coupling vectors $\mathbf{a}$ and $\mathbf{b}$, which add a second component to the update direction and enable geometry-induced momentum, weight decay, and other behaviours.

Table 1: Summary of optimiser variants. All share the same geometric construction (embed, pull back, invert via Sherman-Morrison-Woodbury) and are $O(N)$ per step. Sections 4–6 are the new contributions of this document.

Variant	Base metric γ^{ij}	Learnable?	Section
Fixed diagonal	$\gamma \delta^{ij}$ (hyperparameter)	No	3
Learnable scalar	$e^s \delta^{ij}$	Yes ($s \in \mathbb{R}$)	4
Learnable diagonal	$e^{s_i} \delta^{ij}$	Yes ($\mathbf{s} \in \mathbb{R}^N$)	5
Off-diagonal	$\gamma \delta^{ij} + \text{off-diag. } \mathbf{a}, \mathbf{b}$	No	6

Roadmap. Section 2 establishes notation. Section 3 reproduces the core method of [1], including the full algorithms, and defines the practical machinery (momentum, EMA, log-loss variant) used by all subsequent sections. The three extensions follow in Sections 4–6. Section 7 collects implementation details common to all variants. Appendix A documents the repository structure and sweep infrastructure.

2 Setup and Notation

Given a model with N parameters $\theta(t) = \{\theta^i(t)\}_{i=1}^N$ and a loss function $\mathcal{L} : \mathbb{R}^N \rightarrow \mathbb{R}$, our goal is to find the optimal parameter update $\delta\theta$ at each time step t . That is, in preconditioned gradient flow we have

$$\frac{d\theta^i}{dt} = -g^{ij} \frac{\partial \mathcal{L}}{\partial \theta^j} \implies \delta\theta^i \approx -\eta g^{ij} \frac{\partial \mathcal{L}}{\partial \theta^j} \quad (1)$$

where $\eta \in \mathbb{R}_{>0}$ is the learning rate and g^{ij} is the inverse of the induced metric on the parameter space.

Across all variants in this document the induced metric is obtained by the same geometric construction: we embed parameter space into an ambient space via

$$\phi : \mathbb{R}^N \rightarrow \mathbb{R}^{N+1}, \quad \phi(\theta^1, \dots, \theta^N) = (\theta^1, \dots, \theta^N, f(\theta)),$$

where $f : \mathbb{R}^N \rightarrow \mathbb{R}$ is a scalar function of the parameters (taken to be the loss $\mathcal{L}$, or a function thereof), and then pull back an ambient bilinear form $h \in \mathbb{R}^{(N+1) \times (N+1)}$ to the parameter space. The variants differ in the choice of h (and, for the log-loss variant, the embedding function f ; see below).

We use the shorthand $l_i := \partial\mathcal{L}/\partial\theta^i$ for the raw gradient. When a base metric γ_{ij} is present, we write $l^i := \gamma^{ij}l_j$ for the metric-raised gradient (i.e. the inverse metric applied to $\mathbf{l}$). When $\gamma = \mathbb{I}_N$ (the identity), raising an index has no effect, so $l^i = l_i$ and the two notations coincide. The same index-raising convention applies to any covector: $a^i := \gamma^{ij}a_j$, $m^i := \gamma^{ij}m_j$, etc.

3 Fixed Diagonal Metric

This section reproduces the core method of [1]: a fixed base metric γ defines the ambient space geometry, and the induced metric on parameter space provides a preconditioner for gradient descent. The geometric derivation, practical algorithms, and all notation established here carry forward to the extensions in Sections 4–6.

3.1 Geometric Derivation

The following derivation is due to Thomas [1]; we reproduce it here to establish notation for the extensions that follow.

The ambient metric in [1] is

$$h = \begin{pmatrix} \gamma & \mathbf{0} \\ \mathbf{0}^\top & 1 \end{pmatrix} \in \mathbb{R}^{(N+1) \times (N+1)}$$

where $\gamma \in \mathbb{R}^{N \times N}$ is symmetric positive-definite (taken to be diagonal or scalar in all our experiments). The induced metric g on the parameter space is given by

$$g_{ij} = h_{\alpha\beta} \frac{\partial\phi^\alpha}{\partial\theta^i} \frac{\partial\phi^\beta}{\partial\theta^j} = \gamma_{ij} + \frac{\partial\mathcal{L}}{\partial\theta^i} \frac{\partial\mathcal{L}}{\partial\theta^j}. \quad (2)$$

Then, to calculate the parameter updates $\delta\theta^i$ we need g^{ij} , the inverse of g_{ij} . We can compute this inverse using the Sherman-Morrison-Woodbury formula:

Sherman-Morrison-Woodbury Formula

Let $A \in \mathbb{R}^{n \times n}$ be an invertible matrix, and let $U, V \in \mathbb{R}^{n \times k}$ be matrices. Then, if $C = A + UV^\top$ is invertible, we have

$$C^{-1} = A^{-1} - A^{-1}U(I_k + V^\top A^{-1}U)^{-1}V^\top A^{-1}.$$

This gives us

$$g^{ij} = \gamma^{ij} - \frac{\gamma^{ik}\gamma^{jl} \frac{\partial\mathcal{L}}{\partial\theta^k} \frac{\partial\mathcal{L}}{\partial\theta^l}}{1 + \gamma^{mn} \frac{\partial\mathcal{L}}{\partial\theta^m} \frac{\partial\mathcal{L}}{\partial\theta^n}} = \gamma^{ij} - \frac{l^i l^j}{1 + l_k l^k} \quad (3)$$

where

$$l_i := \frac{\partial\mathcal{L}}{\partial\theta^i} \implies l^i = \gamma^{ij}l_j.$$

Thus, the parameter updates are given by

$$\begin{aligned} \delta\theta^i &= -\eta \left(\gamma^{ij} - \frac{l^i l^j}{1 + l_k l^k} \right) l_j \\ &= -\eta l^i \left(1 - \frac{l_j l^j}{1 + l_k l^k} \right) \\ &= -\frac{\eta l^i}{1 + l_k l^k}. \end{aligned} \quad (4)$$

So to summarize, (i) we start with an ambient metric h which induces a metric g on the parameter space, (ii) we compute the inverse g^{ij} , and (iii) we use g^{ij} to compute the parameter updates $\delta\theta^i$.

For a single component with $\gamma = \mathbb{I}_N$, the update $\delta\theta^i/\eta = -l_i/(1 + \sum_k l_k^2)$ depends on both l_i and the other gradients $\{l_k\}_{k \neq i}$. Figure 1 plots this as a function of l_i for several fixed values of $\sum_{k \neq i} l_k^2$.




Figure 1: Step-size profile for the fixed diagonal metric ($\gamma = \mathbb{I}$, $\xi = 1$): $\delta\theta^i/\eta = -l_i/(1 + \sum_k l_k^2)$ plotted as a function of l_i for several fixed values of $\sum_{k \neq i} l_k^2$. For small gradients the update is approximately linear (like SGD); for large gradients it saturates, providing smooth gradient clipping. More background curvature from the other parameters produces stronger damping. Compare Figure 2 of [1].

3.2 Practical Algorithm

The theoretical update (4) is clean but not directly usable in stochastic training. Thomas [1] bridges this gap with four modifications, each of which is standard in the optimisation literature but takes on a specific role in the induced-metric framework.

Metric strength ξ . The theoretical derivation implicitly sets $\xi = 1$. Geometrically, ξ enters through the embedding itself: one replaces $\phi(\theta) = (\theta, f(\theta))$ with $\phi(\theta) = (\theta, \sqrt{\xi} f(\theta))$, which scales the loss dimension of the ambient space and yields the induced metric $g_{ij} = \gamma_{ij} + \xi l_i l_j$. The parameter $\xi > 0$ therefore controls how strongly the loss-surface curvature influences the preconditioner: $\xi \rightarrow 0$ recovers the base metric γ^{-1} , while large ξ makes the geometric clipping dominant.

Momentum. An exponential moving average (EMA) of the raw gradients,

$$\mathbf{m}_t = \beta_m \mathbf{m}_{t-1} + (1 - \beta_m) \mathbf{l}_t, \quad \mathbf{m}_0 = \mathbf{0}, \quad (5)$$

replaces the instantaneous gradient with a smoothed direction. Here $\beta_m \in [0, 1)$ is the momentum decay: $\beta_m = 0$ recovers the raw gradient, while β_m close to 1 gives a long look-back (roughly $1/(1 - \beta_m)$ steps). The bias-corrected, metric-raised momentum $\hat{\mathbf{m}}_t = \gamma^{-1} \mathbf{m}_t / (1 - \beta_m^t)$ is then used as the step direction. Momentum is not required by the geometric framework—the induced metric works without it—but it provides the standard acceleration benefits (variance reduction, faster traversal of flat regions) and is included to match common practice.

Denominator EMA. The scalar $v_t = \xi \mathbf{l}_t^\top \gamma^{-1} \mathbf{l}_t$ is the metric-weighted squared gradient norm. Thomas [1] calls it the *trace statistic* because it equals $\xi \cdot \text{tr}(\gamma^{-1} \mathbf{l} \mathbf{l}^\top)$, the trace of the rank-1 correction to the induced metric. In stochastic training each minibatch (a random subset of the data) yields a different v_t . To

stabilise the damping factor, v_t is smoothed with a *second* EMA, independent of the momentum EMA above (momentum smooths the gradient *direction*; this one smooths the scalar *damping*):

$$\bar{v}_t = \beta \bar{v}_{t-1} + (1 - \beta) v_t, \quad \bar{v}_0 = 0, \quad (6)$$

bias-corrected to $\hat{v}_t = \bar{v}_t / (1 - \beta^t)$. The parameter $\beta \in [0, 1)$ plays the same role as β_m : it sets the look-back window, roughly $1/(1 - \beta)$ steps. The damping factor is then $r_t = 1/(1 + |\hat{v}_t|)$.

Decoupled weight decay. Following [1], weight decay is applied directly to the parameters ($+\lambda\theta_{t-1}$ with $\lambda \leq 0$ for decay) rather than through the loss gradient. This decoupled form is the natural choice from the geometric perspective due to general covariance [1]. Like momentum, weight decay is not required by the induced-metric framework but is included because it is standard practice and has a natural geometric justification.

These modifications are collected in Algorithm 1 (adapted from Algorithm 1 of [1]). The log-loss variant is Algorithm 2 (adapted from Algorithm 2 of [1]).

Algorithm 1 Fixed diagonal metric (plain embedding, $f = \mathcal{L}$)

Require: learning rate η , momentum decay β_m , metric strength ξ , denominator EMA decay β , weight decay λ , inverse base metric γ^{-1}

```

1: Initialise  $t \leftarrow 0$ ,  $\mathbf{m}_0 \leftarrow \mathbf{0}$ ,  $\bar{v}_0 \leftarrow 0$ 
2: while  $\theta_t$  not converged do
3:    $t \leftarrow t + 1$ 
4:    $\mathbf{l}_t \leftarrow \nabla_{\theta} \mathcal{L}(\theta_{t-1})$ 
5:    $\mathbf{v}_t \leftarrow \xi \cdot \sum_i \mathbf{l}_{t,i} (\gamma^{-1} \mathbf{l}_t)_i$  // trace statistic
6:    $\bar{\mathbf{v}}_t \leftarrow \beta \bar{\mathbf{v}}_{t-1} + (1 - \beta) \mathbf{v}_t$  // denominator EMA
7:    $\hat{\mathbf{v}}_t \leftarrow \bar{\mathbf{v}}_t / (1 - \beta^t)$  // bias correction
8:    $r_t \leftarrow 1/(1 + |\hat{\mathbf{v}}_t|)$  // damping factor
9:    $\mathbf{m}_t \leftarrow \beta_m \mathbf{m}_{t-1} + (1 - \beta_m) \mathbf{l}_t$  // momentum
10:   $\hat{\mathbf{m}}_t \leftarrow \gamma^{-1} \mathbf{m}_t / (1 - \beta_m^t)$  // bias-corrected, metric-raised
11:   $\theta_t \leftarrow \theta_{t-1} - \eta r_t \hat{\mathbf{m}}_t + \lambda \theta_{t-1}$ 
12: end while
13: return  $\theta_t$ 
```

Algorithm 2 Fixed diagonal metric (log-loss embedding, $f = \log \mathcal{L}$, requires $\mathcal{L} > 0$)

Require: same as Algorithm 1

```

1: Initialise  $t \leftarrow 0$ ,  $\mathbf{m}_0 \leftarrow \mathbf{0}$ ,  $\bar{v}_0 \leftarrow 0$ 
2: while  $\theta_t$  not converged do
3:    $t \leftarrow t + 1$ 
4:    $\mathbf{l}_t \leftarrow \nabla_{\theta} \mathcal{L}(\theta_{t-1})$ 
5:    $L_t \leftarrow \mathcal{L}(\theta_{t-1})$ 
6:    $\mathbf{v}_t \leftarrow \xi \cdot \sum_i \mathbf{l}_{t,i} (\gamma^{-1} \mathbf{l}_t)_i$ 
7:    $\bar{\mathbf{v}}_t \leftarrow \beta \bar{\mathbf{v}}_{t-1} + (1 - \beta) \mathbf{v}_t$ 
8:    $\hat{\mathbf{v}}_t \leftarrow \bar{\mathbf{v}}_t / (1 - \beta^t)$ 
9:    $r_t \leftarrow L_t / (L_t^2 + |\hat{\mathbf{v}}_t|)$  // log-loss damping
10:   $\mathbf{m}_t \leftarrow \beta_m \mathbf{m}_{t-1} + (1 - \beta_m) \mathbf{l}_t$ 
11:   $\hat{\mathbf{m}}_t \leftarrow \gamma^{-1} \mathbf{m}_t / (1 - \beta_m^t)$ 
12:   $\theta_t \leftarrow \theta_{t-1} - \eta r_t \hat{\mathbf{m}}_t + \lambda \theta_{t-1}$ 
13: end while
14: return  $\theta_t$ 
```

In short, the practical update at each step is

$$\delta\theta^i = -\eta r_t \hat{m}^i, \quad r_t = \frac{1}{1 + |\hat{v}_t|}, \quad v_t = \xi \gamma^{mn} l_{t,m} l_{t,n}, \quad (7)$$

where $\hat{v}_t$ is the bias-corrected EMA of v_t (lines 5–8 of Algorithm 1) and $\hat{m}^i$ is the bias-corrected, metric-raised momentum (line 10).

3.3 Log-Loss Variant

Log-Loss Derivation Sketch

With $f = \log \mathcal{L}$ the embedding gradient is $\partial f / \partial \theta^i = l_i / \mathcal{L}$, so the induced metric becomes $g_{ij} = \gamma_{ij} + \xi (l_i / \mathcal{L})(l_j / \mathcal{L})$. Applying Sherman-Morrison-Woodbury gives the inverse metric

$$g^{ij} = \gamma^{ij} - \frac{\xi l^i l^j}{\mathcal{L}^2 + \xi l_k l^k}.$$

Contracting with the *raw* gradient l_j yields $g^{ij} l_j = l^i \cdot \mathcal{L}^2 / (\mathcal{L}^2 + v_t)$, where $v_t = \xi \gamma^{mn} l_m l_n$.

The full geometric update, however, uses the embedding gradient $\partial_j f = l_j / \mathcal{L}$, not the raw gradient:

$$g^{ij} \frac{\partial f}{\partial \theta^j} = \frac{g^{ij} l_j}{\mathcal{L}} = l^i \cdot \frac{\mathcal{L}}{\mathcal{L}^2 + v_t}.$$

The practical damping factor $r_t = \mathcal{L}_t / (\mathcal{L}_t^2 + |\hat{v}_t|)$ (line 7 of Algorithm 2) therefore follows directly from the geometry of the log embedding. This provides stronger damping than $\mathcal{L}^2 / (\mathcal{L}^2 + |\hat{v}_t|)$ as $\mathcal{L} \rightarrow 0$, preventing overshooting near minima.

In practice, the update direction $\mathbf{r}$ and the metric direction $\mathbf{l}$ need not coincide; see Section 7 for the full treatment.

3.4 Implementation

- JAX/Optax: `optimisers/jax_fixed.py`
- PyTorch: `optimisers/torch_fixed.py`
- Equivalence check: `optimisers/compare_fixed.py`

Exposed optimisers: `custom_sgd` (plain embedding, $f = \mathcal{L}$), `custom_sgd_log` (log-loss embedding, $f = \log \mathcal{L}$; requires $\mathcal{L} > 0$), `custom_sgd_rms` (RMS-normalised gradients, i.e. the update direction is $l_i / \|\mathbf{l}\|$ instead of l_i). The sweep name is `sgd_metric` / `sgd_log_metric` / `sgd_rms`.

4 Learnable Scalar Metric

The fixed diagonal variant treats γ as a hyperparameter chosen before training. A natural next step is to let the metric adapt during optimisation. We begin with the simplest case: a single learnable scalar that controls the global scale of γ^{-1} .

4.1 Construction

We parameterise the inverse metric as

$$\gamma^{ij} = \alpha \delta^{ij}, \quad \alpha = e^s,$$

where $s \in \mathbb{R}$ is a single learnable scalar. (We learn the *inverse* metric directly because the update rule requires γ^{-1} .) Initialising $s = 0$ gives $\gamma^{-1} = \mathbb{1}_N$, so the optimiser starts as the unscaled fixed diagonal variant.

The ambient metric remains block-diagonal,

$$h = \begin{pmatrix} e^{-s} \mathbb{1}_N & \mathbf{0} \\ \mathbf{0}^\top & 1 \end{pmatrix},$$

and the induced metric on parameter space is

$$g_{ij} = e^{-s} \delta_{ij} + l_i l_j.$$

Applying the Sherman-Morrison-Woodbury formula with $A_{ij} = e^{-s} \delta_{ij}$, $u = v = \mathbf{l}$:

$$g^{ij} = \gamma^{ij} - \frac{\gamma^{ik} l_k \gamma^{jl} l_l}{1 + \gamma^{mn} l_m l_n} = e^s \delta^{ij} - \frac{e^{2s} l_i l_j}{1 + e^s \|\mathbf{l}\|^2}. \quad (8)$$

The parameter update is therefore

$$\delta \theta^i = -\eta g^{ij} l_j = -\frac{\eta \gamma^{ij} l_j}{1 + \gamma^{mn} l_m l_n} = -\frac{\eta e^s l_i}{1 + e^s \|\mathbf{l}\|^2}. \quad (9)$$

The denominator depends on s through $\gamma^{mn} = e^s \delta^{mn}$. Figure 2 shows how varying s reshapes the step-size profile relative to the fixed diagonal case (Figure 1).

The practical update follows the same structure as (7), with $\gamma^{ij} = \alpha \delta^{ij}$ and the trace statistic specialised to

$$v_t = \xi \alpha \|\mathbf{l}_t\|^2. \quad (10)$$

4.2 Metric Learning

We want to learn s so that $\alpha = e^s$ is calibrated to the local geometry of the loss. The trace statistic $v_t = \xi \alpha \|\mathbf{l}_t\|^2$ appears directly in the denominator of the update rule through $r_t = 1/(1 + |\hat{v}_t|)$. When v_t is near zero the damping is inactive and the optimiser reduces to scaled SGD; when v_t is large the damping dominates and the geometric clipping is fully engaged. We adapt s by performing *gradient ascent* on v_t with respect to s . This is a convenient heuristic— v_t is the quantity that already appears in the denominator, and ascent on it has a clear geometric interpretation (see below)—but we do not claim it is the unique or optimal choice of surrogate. Concretely:

$$\frac{\partial v_t}{\partial s} = \xi e^s \|\mathbf{l}_t\|^2 = v_t.$$

Why ascent? Consider a single parameter with $\gamma = 1$. The effective step size per unit gradient is

$$\frac{|\delta \theta|}{|l|} = \frac{e^s}{1 + \xi e^s l^2}. \quad (11)$$

Its derivative with respect to the metric parameter s is

$$\frac{\partial}{\partial s} \left(\frac{e^s}{1 + \xi e^s l^2} \right) = \frac{e^s}{(1 + \xi e^s l^2)^2} > 0. \quad (12)$$

The step size is therefore *monotonically increasing* in s (note: increasing in s , not in l), bounded above by $1/(\xi l^2)$. Gradient ascent on v increases s , pushing the step size toward this upper bound—the regime where the geometric clipping is most active. Gradient *descent* would decrease s , collapsing the step size toward zero and disabling the metric entirely. The regularisation term (below) provides a restoring force, and the equilibrium between ascent and regularisation determines the operating point.




Figure 2: Step-size profile for the learnable scalar metric: $\delta\theta^i/\eta = -e^s l_i / (1 + \xi e^s l_i^2)$ in the single-parameter case ($N = 1$, $\xi = 1$). Increasing s amplifies the peak step size and shifts it toward smaller gradients; decreasing s flattens the curve, approaching plain SGD. The dashed line is the fixed diagonal baseline ($s = 0$).

Regularisation and stability. Since $v_t \geq 0$, the unregularised ascent rule monotonically increases s . Three mechanisms keep s bounded: (i) an L2 regularisation term $-\mu\lambda_s s$ that pulls s back towards zero (i.e. $\gamma^{-1} = \mathbb{1}$), (ii) hard clipping of s to $[-s_{\max}, s_{\max}]$, and (iii) the denominator EMA, which smooths transient gradient spikes.

The online metric update is

$$s \leftarrow s + \mu v_t - \mu \lambda_s s, \quad (13)$$

where $\mu > 0$ is the metric learning rate and $\lambda_s > 0$ is the L2 regularisation coefficient, followed by clipping to $[-s_{\max}, s_{\max}]$.




Figure 3: Evolution of the learnable scalar s during training (2-layer MLP on synthetic regression, 500 steps). s rises rapidly in early training as the metric engages with the gradient structure, then plateaus as the regularisation balances the ascent pressure. The loss (right axis, log scale) decreases throughout.

4.3 Algorithm

Algorithm 3 Learnable scalar metric (plain embedding)

Require: all hyperparameters from Algorithm 1, plus metric learning rate μ , metric regularisation λ_s , clipping bound $s_{\max}$

```

1: Initialise  $t \leftarrow 0$ ,  $\mathbf{m}_0 \leftarrow \mathbf{0}$ ,  $\bar{v}_0 \leftarrow 0$ ,  $s \leftarrow 0$ 
2: while  $\theta_t$  not converged do
3:    $t \leftarrow t + 1$ 
4:    $\alpha \leftarrow e^s$ 
5:    $\mathbf{l}_t \leftarrow \nabla_{\theta} \mathcal{L}(\theta_{t-1})$ 
6:    $v_t \leftarrow \xi \alpha \|\mathbf{l}_t\|^2$ 
7:    $\bar{v}_t \leftarrow \beta \bar{v}_{t-1} + (1 - \beta) v_t$ 
8:    $\hat{v}_t \leftarrow \bar{v}_t / (1 - \beta^t)$ 
9:    $r_t \leftarrow 1 / (1 + |\hat{v}_t|)$ 
10:   $\mathbf{m}_t \leftarrow \beta_m \mathbf{m}_{t-1} + (1 - \beta_m) \mathbf{l}_t$ 
11:   $\hat{\mathbf{m}}_t \leftarrow \alpha \mathbf{m}_t / (1 - \beta_m^t)$ 
12:   $\theta_t \leftarrow \theta_{t-1} - \eta r_t \hat{\mathbf{m}}_t + \lambda \theta_{t-1}$  // parameter update
13:   $s \leftarrow s + \mu v_t - \mu \lambda_s s$  // metric update
14:   $s \leftarrow \text{clip}(s, -s_{\max}, s_{\max})$ 
15: end while
16: return  $\theta_t$ 

```

The log-loss variant replaces line 6 with $r_t \leftarrow L_t / (L_t^2 + |\hat{v}_t|)$.

4.4 Hyperparameters

The learnable scalar variant introduces three new hyperparameters beyond those of the fixed diagonal variant:

- μ (`metric_lr`): step size for the online metric update in (13).
- λ_s (`metric_reg`): L2 regularisation strength on s , controlling how far the metric drifts from identity.
- $s_{\max}$ (`metric_clip`): clipping bound on s , bounding $\alpha \in [e^{-s_{\max}}, e^{s_{\max}}]$.

4.5 Implementation

- JAX/Optax: `optimisers/jax_learnable_scalar.py`
- PyTorch: `optimisers/torch_learnable_scalar.py`
- Equivalence check: `optimisers/compare_learnable_scalar.py`

The plain variant is `custom_sgd_learnable_scalar` (sweep: `sgd_learn_scalar`). The log-loss variant is `custom_sgd_log_learnable_scalar` (sweep: `sgd_learn_scalar_log`).

5 Learnable Diagonal Metric

The learnable scalar metric has a single degree of freedom and therefore applies the same scale to every parameter. We now generalise to a per-parameter learnable scale, giving each parameter its own adaptive metric entry.

5.1 Construction

We promote s to a full vector $\mathbf{s} \in \mathbb{R}^N$ with one entry per parameter, giving the per-parameter inverse metric

$$\gamma^{ij} = e^{s_i} \delta^{ij}.$$

The ambient metric is

$$h = \begin{pmatrix} \text{diag}(e^{-\mathbf{s}}) & \mathbf{0} \\ \mathbf{0}^\top & 1 \end{pmatrix},$$

and the induced metric on parameter space is

$$g_{ij} = e^{-s_i} \delta_{ij} + l_i l_j.$$

This is again a rank-1 update to a diagonal matrix. Applying Sherman-Morrison-Woodbury with $A_{ij} = e^{-s_i} \delta_{ij}$ and $u = v = \mathbf{l}$,

$$g^{ij} = e^{s_i} \delta^{ij} - \frac{l^i l^j}{1 + l_k l^k},$$

where $l^i = \gamma^{ij} l_j = e^{s_i} l_i$ (no sum on i), so $l^i l^j = e^{s_i} l_i \cdot e^{s_j} l_j$ and $l_k l^k = \sum_k e^{s_k} l_k^2$. The update is

$$\delta \theta^i = -\eta g^{ij} l_j = -\eta \left(l^i - \frac{l^i (l_k l^k)}{1 + l_k l^k} \right) = -\frac{\eta l^i}{1 + l_k l^k}. \quad (14)$$

With ξ , momentum, and the EMA damping, the practical update is

$$\delta \theta^i = -\eta r_t \hat{m}^i, \quad r_t = \frac{1}{1 + |\hat{v}_t|}, \quad v_t = \xi \sum_k e^{s_k} l_k^2, \quad (15)$$

with the same log-loss variant as before. Compared to the scalar case, each parameter θ^i now has an individual scale e^{s_i} , allowing the metric to adapt separately to different regions of parameter space.

5.2 Metric Learning

The per-parameter metric update treats each s_i independently. We use the same surrogate-ascent strategy as in Section 4.2, now applied to the trace statistic $v_t = \xi \sum_k e^{s_k} l_k^2$ which serves as the surrogate objective:

$$J(\mathbf{s}) = v_t = \xi \sum_k e^{s_k} l_k^2.$$

This is the quantity that appears in the denominator of the update rule, so adapting $\mathbf{s}$ to calibrate v_t keeps the damping in a useful operating range (see Section 4.2 for the detailed argument). Its gradient with respect to s_i is

$$\frac{\partial J}{\partial s_i} = \xi e^{s_i} l_i^2.$$

The online update is therefore

$$s_i \leftarrow s_i + \mu \xi e^{s_i} l_i^2 - \mu \lambda_s s_i, \quad (16)$$

followed by mean-centering and clipping (discussed below). The motivation for gradient *ascent* (not descent) and the role of regularisation are the same as in Section 4.2.

5.3 Mean-Centering

Parameter tensors and leaves. A model’s parameters are organised as a collection of tensors (e.g. one weight matrix and one bias vector per layer). Each such tensor is called a *leaf*. We write N_ℓ for the number of scalar entries in a single leaf, as distinct from the total parameter count $N = \sum_\ell N_\ell$.

The centering operation. The vector $\mathbf{s}$ has one entry per parameter, $i = 1, \dots, N$. After the gradient step (16), we mean-centre $\mathbf{s}$ *within each leaf independently*: for each leaf ℓ containing parameters $i \in I_\ell$ (with $|I_\ell| = N_\ell$),

$$s_i \leftarrow s_i - \bar{s}_\ell, \quad \bar{s}_\ell = \frac{1}{N_\ell} \sum_{k \in I_\ell} s_k.$$

This is not an ad-hoc trick but a principled normalisation arising from a gauge degeneracy.

Gauge degeneracy. Write each log-scale as a mean plus a deviation: $s_k = \bar{s} + \delta_k$. By construction $\sum_k \delta_k = \sum_k (s_k - \bar{s}) = N_\ell \bar{s} - N_\ell \bar{s} = 0$. The trace statistic then factorises as

$$v = \xi \sum_k e^{s_k} l_k^2 = \xi e^{\bar{s}} \sum_k e^{\delta_k} l_k^2.$$

The global mean $\bar{s}$ acts as a multiplicative prefactor that is exactly degenerate with ξ : shifting all s_k by a constant c is equivalent to scaling ξ by e^c and leaving $\mathbf{s}$ unchanged. Mean-centering enforces $\bar{s} = 0$ (note: subsequent clipping may re-introduce a small nonzero mean), which fixes this gauge and ensures that ξ remains the single scalar controlling the overall metric magnitude while $\mathbf{s}$ captures only the *relative* per-parameter geometry. Since we apply this per leaf rather than globally, a separate mean mode is removed in each parameter tensor.

Fixed Diagonal as a Special Case

In a highly symmetric special case: if all s_i are initialised to the same value and the gradients have the same magnitude for all parameters at every step, then $\mathbf{s}$ stays uniform and mean-centering keeps it at zero. In that case $e^{s_i} = 1$ for all i and $v = \xi \|\mathbf{l}\|^2$, which reduces exactly to the fixed diagonal variant with $\gamma = \mathbb{1}_N$. The learnable diagonal metric therefore generalises fixed diagonal by learning *relative* per-parameter scalings, while mean-centering intentionally removes the global scale mode

(which is degenerate with ξ). To recover a learnable global scale, disable mean-centering or make ξ itself learnable.




Figure 4: Evolution of per-parameter log-scales s_i during training (same setup as Figure 3). Different parameters acquire different scales: some are upscaled ($s_i > 0$), others downscaled ($s_i < 0$). Mean-centering keeps the average at zero, so $\mathbf{s}$ captures only relative geometry.

5.4 Algorithm

Algorithm 4 Learnable diagonal metric (plain embedding)

Require: all hyperparameters from Algorithm 1, plus μ , λ_s , $s_{\max}$

```

1: Initialise  $t \leftarrow 0$ ,  $\mathbf{m}_0 \leftarrow \mathbf{0}$ ,  $\bar{v}_0 \leftarrow 0$ ,  $\mathbf{s} \leftarrow \mathbf{0}$ 
2: while  $\theta_t$  not converged do
3:    $t \leftarrow t + 1$ 
4:    $\mathbf{l}_t \leftarrow \nabla_{\theta} \mathcal{L}(\theta_{t-1})$ 
5:    $\mathbf{v}_t \leftarrow \xi \sum_k e^{s_k} l_{t,k}^2$ 
6:    $\bar{v}_t \leftarrow \beta \bar{v}_{t-1} + (1 - \beta) v_t$ 
7:    $\hat{v}_t \leftarrow \bar{v}_t / (1 - \beta^t)$ 
8:    $r_t \leftarrow 1 / (1 + |\hat{v}_t|)$ 
9:    $\mathbf{m}_t \leftarrow \beta_m \mathbf{m}_{t-1} + (1 - \beta_m) \mathbf{l}_t$ 
10:   $\hat{\mathbf{m}}_t^i \leftarrow e^{s_i} m_{t,i} / (1 - \beta_m^t)$ 
11:   $\theta_t \leftarrow \theta_{t-1} - \eta r_t \hat{\mathbf{m}}_t + \lambda \theta_{t-1}$ 
12:   $s_i \leftarrow s_i + \mu \xi e^{s_i} l_{t,i}^2 - \mu \lambda_s s_i$  for each  $i$ 
13:   $s_i \leftarrow s_i - \bar{s}$  per leaf
14:   $\mathbf{s} \leftarrow \text{clip}(\mathbf{s}, -s_{\max}, s_{\max})$ 
15: end while
16: return  $\theta_t$ 

```

// parameter update

// metric gradient ascent

// mean-centre (gauge fix)

5.5 Implementation

- JAX/Optax: `optimisers/jax_learnable_diag.py`
- PyTorch: `optimisers/torch_learnable_diag.py`
- Equivalence check: `optimisers/compare_learnable_diag.py`

The plain variant is `custom_sgd_learnable_diag` (sweep: `sgd_learn_diag`). The log-loss variant is `custom_sgd_log_learnable_diag` (sweep: `sgd_learn_diag_log`).

6 Off-Diagonal Metric

All preceding variants use an ambient metric h whose off-diagonal blocks (between the N parameter dimensions and the single loss dimension) are zero. We now explore what happens when those blocks are nonzero—a direction explicitly suggested by [1].

6.1 Derivation

Consider a more general ambient bilinear form

$$h = \begin{pmatrix} \gamma & \mathbf{a} \\ \mathbf{b}^\top & 1 \end{pmatrix} \in \mathbb{R}^{(N+1) \times (N+1)}$$

where $\mathbf{a}, \mathbf{b} \in \mathbb{R}^N$. Intuitively, the vectors $\mathbf{a}$ and $\mathbf{b}$ introduce off-diagonal couplings between the parameter space and the loss function dimension. In general, if $\mathbf{a} \neq \mathbf{b}$ this h is not symmetric, so here we simply treat it as a preconditioning matrix rather than a Riemannian metric.

Stability note. If the effective preconditioner is non-symmetric ($\mathbf{a} \neq \mathbf{b}$), the update direction need not be a descent direction even for convex objectives. To see why, note that a preconditioned gradient step $\delta\theta = -\eta P \nabla \mathcal{L}$ is a descent direction only if $\nabla \mathcal{L}^\top P \nabla \mathcal{L} > 0$, which is guaranteed when P is positive-definite (or has positive-definite symmetric part). A non-symmetric h can produce a non-symmetric $P = g^{-1}$, for which the quadratic form $\nabla \mathcal{L}^\top P \nabla \mathcal{L}$ can be negative—meaning the update moves *uphill*. We therefore recommend defaulting to the symmetric case $\mathbf{a} = \mathbf{b}$; if using $\mathbf{a} \neq \mathbf{b}$, consider symmetrising the effective preconditioner or otherwise enforcing positive-definiteness.

The induced bilinear form on the parameter space is

$$\begin{aligned} g_{ij} &= h_{\alpha\beta} \frac{\partial \phi^\alpha}{\partial \theta^i} \frac{\partial \phi^\beta}{\partial \theta^j} \\ &= \gamma_{ij} + \xi(a_i l_j + b_j l_i + l_i l_j). \end{aligned}$$

To compute the inverse g^{ij} , we notice that g can be expressed as a rank-2 update to γ : for $U = [\mathbf{a}, \mathbf{l}]$, $V = [\mathbf{l}, \mathbf{b} + \mathbf{l}] \in \mathbb{R}^{N \times 2}$ we have

$$\begin{aligned} \gamma_{ij} + \xi U_{ik} V_{jk} &= \gamma_{ij} + \xi(a_i l_j + (b_j + l_j) l_i) \\ &= \gamma_{ij} + \xi(a_i l_j + b_j l_i + l_i l_j) \\ &= g_{ij}. \end{aligned}$$

Applying the Sherman-Morrison-Woodbury formula to $g = \gamma + \xi UV^\top$:

$$g^{ij} = \gamma^{ij} - \xi \gamma^{ik} U_{k\alpha} \left(\delta_{\alpha\beta} + \xi V_{m\alpha} \gamma^{mn} U_{n\beta} \right)^{-1} V_{l\beta} \gamma^{lj}.$$

For convenience we define the scalar contractions

$$c_{al} := a_n l^n, \quad c_{ul} := l_n l^n, \quad c_{ba} := b_n a^n, \quad c_{bl} := b_n l^n,$$

so that the 2×2 core matrix is

$$\mathbb{1}_2 + \xi V^\top \gamma^{-1} U = \begin{pmatrix} 1 + \xi c_{al} & \xi c_{ul} \\ \xi(c_{ba} + c_{al}) & 1 + \xi(c_{bl} + c_{ul}) \end{pmatrix}$$

with determinant

$$D = 1 + \xi(c_{al} + c_{bl} + c_{ul}) + \xi^2(c_{al} c_{bl} - c_{ul} c_{ba}). \quad (17)$$

Inverting the 2×2 system and contracting yields the parameter update

$$\delta\theta^i = -\eta g^{ij} l_j = -\eta \left[\frac{1 + \xi c_{al}}{D} l^i - \frac{\xi c_{ul}}{D} a^i \right]. \quad (18)$$

The closed-form inverse requires $D \neq 0$ (equivalently, $\det(\mathbb{1}_2 + \xi V^\top \gamma^{-1} U) \neq 0$). The JAX implementation includes a near-singular fallback; the PyTorch implementation solves the 2×2 system directly via `torch.linalg.solve`.

A numerical check of this has been performed in `documents/off-diag-check.ipynb`: for random $\mathbf{a}, \mathbf{b}, \gamma, \mathbf{l}$ the closed-form $g^{ij} l_j$ matches `torch.linalg.solve(g, l)` within floating-point tolerance.

Sanity Checks

Diagonal recovery ($\mathbf{a} = \mathbf{b} = \mathbf{0}$): $c_{al} = c_{ba} = c_{bl} = 0$, so $D = 1 + \xi c_{ul}$. Equation (18) gives $\delta\theta^i = -\eta l^i / (1 + \xi c_{ul})$, which matches the fixed diagonal case. ✓

Symmetric coupling ($\mathbf{a} = \mathbf{b} = \mathbf{l}$): $c_{al} = c_{bl} = c_{ba} = c_{ul}$, so $D = 1 + 3\xi c_{ul} + \xi^2(c_{ul}^2 - c_{ul}^2) = 1 + 3\xi c_{ul}$. The update simplifies to $\delta\theta^i = -\eta l^i / (1 + 3\xi c_{ul})$. ✓

One-sided coupling ($\mathbf{a} = \mathbf{0}, \mathbf{b} \neq \mathbf{0}$): $c_{al} = c_{ba} = 0$, $D = 1 + \xi(c_{bl} + c_{ul})$. The update reduces to $\delta\theta^i = -\eta l^i / D$, i.e. a scaled gradient step whose denominator depends on $\mathbf{b}$ through c_{bl} .

The update direction (18) has two components: one along l^i (the gradient direction) and one along a^i (the off-diagonal coupling direction). Depending on the choices of $\mathbf{a}$ and $\mathbf{b}$, this can lead to a variety of behaviours:

- If $\mathbf{a} = \mathbf{0}$ or $\mathbf{a} \propto \mathbf{l}$: plain gradient descent with an adaptive learning rate.
- If $\mathbf{a} = \text{momentum}$: a momentum-like update with an adaptive learning rate.
- If $\mathbf{a} = \boldsymbol{\theta}$ (the parameter vector): gradient descent combined with a form of geometry-induced weight decay or growth, depending on the sign of $\xi c_{ul}/D$.




Figure 5: Step-size profiles for off-diagonal variants in the single-parameter case ($N = 1$, $\gamma = 1$, $\xi = 1$). All curves plot $\delta\theta/\eta = -l/(1 + \xi_{\text{eff}} l^2)$ where ξ_{eff} depends on the coupling: $\xi_{\text{eff}} = 1$ (fixed diagonal, $\mathbf{a} = \mathbf{b} = \mathbf{0}$), $\xi_{\text{eff}} = 2$ (one-sided), $\xi_{\text{eff}} = 3$ (symmetric, $\mathbf{a} = \mathbf{b} = \mathbf{l}$).

6.2 Algorithm

Algorithm 5 Off-diagonal metric

Require: learning rate η , momentum decay β_m , metric strength ξ , weight decay λ , inverse base metric γ^{-1} , mode choices for $\mathbf{a}$, $\mathbf{b}$, $\mathbf{l}$, $\mathbf{r}$

- 1: Initialise $t \leftarrow 0$, $\mathbf{m}_0 \leftarrow \mathbf{0}$
 - 2: **while** θ_t not converged **do**
 - 3: $t \leftarrow t + 1$
 - 4: $\mathbf{g}_t \leftarrow \nabla_{\theta} \mathcal{L}(\theta_{t-1})$
 - 5: $\mathbf{m}_t \leftarrow \beta_m \mathbf{m}_{t-1} + (1 - \beta_m) \mathbf{g}_t$; $\hat{\mathbf{m}}_t \leftarrow \mathbf{m}_t / (1 - \beta_m^t)$
 - 6: Select $\mathbf{a}$, $\mathbf{b}$, $\mathbf{l}$, $\mathbf{r}$ from modes (gradient, momentum, θ , or $\mathbf{0}$)
 - 7: $\mathbf{U} \leftarrow (\mathbf{a}, \mathbf{l})$; $\mathbf{V} \leftarrow (\mathbf{l}, \mathbf{b} + \mathbf{l})$
 - 8: Compute c_{al} , c_{ll} , c_{ba} , c_{bl} and D via (17)
 - 9: Compute $G^{-1} \mathbf{r}$ via Woodbury (2×2 solve + $O(N)$ vector ops)
 - 10: $\theta_t \leftarrow \theta_{t-1} - \eta (G^{-1} \mathbf{r}) + \lambda \theta_{t-1}$
 - 11: **end while**
 - 12: **return** θ_t
-

Note that Algorithm 5 does not use a denominator EMA (β): the full 2×2 Woodbury solve is applied directly at each step.

6.3 Implementation

- JAX/Optax: `optimisers/jax_offdiag.py`
- PyTorch: `optimisers/torch_offdiag.py`

- Equivalence check: `optimisers/compare_offdiag.py`

Sweep names follow the pattern `sgd_offdiag_{a}_{b}` where `{a}` and `{b}` are drawn from 0 (zero), 1 (gradient), `m` (momentum), `theta` (parameters). All pairs except `sgd_offdiag_0_0` (which reduces to the fixed diagonal variant) are registered in `parameters/optimizer_registry.py`.

Computational cost. Taking $\gamma_{ij} = \gamma \delta_{ij}$ (scalar base metric) and writing

$$G = \gamma \mathbb{1}_N + \xi UV^\top, \quad U = (\mathbf{a}, \mathbf{l}), \quad V = (\mathbf{l}, \mathbf{b} + \mathbf{l}),$$

the Sherman-Morrison-Woodbury formula gives

$$G^{-1}\mathbf{r} = \frac{1}{\gamma}\mathbf{r} - \frac{\xi}{\gamma^2}U\left(\mathbb{1}_2 + \frac{\xi}{\gamma}V^\top U\right)^{-1}V^\top\mathbf{r}.$$

This costs $O(N)$: only inner products and a single 2×2 matrix inversion are required, never a full $N \times N$ matrix.

Modes for $\mathbf{l}$, $\mathbf{a}$, and $\mathbf{b}$. The vector $\mathbf{l}$ encodes the direction along which the loss coordinate stretches, with two built-in choices:

- **Gradient mode:** $\mathbf{l}_i = \partial_i \mathcal{L}$ — the metric is shaped by the instantaneous slope.
- **Momentum mode:** $\mathbf{l}_i = \mathbf{m}_i$ — the metric incorporates accumulated gradient history.

The off-diagonal coupling vectors $\mathbf{a}$ and $\mathbf{b}$ each support:

- **Zero:** $\mathbf{a}_i = 0$ or $\mathbf{b}_i = 0$ — eliminates one branch of the coupling.
- **Gradient:** aligned with $\partial_i \mathcal{L}$.
- **Momentum:** aligned with $\mathbf{m}_i$.
- **Parameter:** $\mathbf{a}_i = \theta^i$ or $\mathbf{b}_i = \theta^i$ — couples geometry to position.
- **Same-as-a:** $\mathbf{b}_i = \mathbf{a}_i$ — symmetric off-diagonal structure.

For full generality the implementations also accept user-provided vectors or callables.

7 Shared Implementation Details

7.1 Update Direction vs. Metric Direction

In all implementations, the goal is to realise

$$\delta\theta^i = -\eta g^{ij}r_j,$$

where r_j is the *update direction* and $\mathbf{l}_i$ is the direction used to *construct the metric*. These need not be the same vector. (Note that r_j here is a vector, not to be confused with the scalar damping factor r_t in Sections 3–5; the distinction is always clear from context and index type.) In ordinary gradient descent one would take $r_j = \mathbf{l}_j = \partial_j \mathcal{L}$, but we also test momentum-based choices. The four combinations are:

- $\mathbf{l} = \nabla \mathcal{L}$, $\mathbf{r} = \nabla \mathcal{L}$: pure geometric gradient descent.
- $\mathbf{l} = \nabla \mathcal{L}$, $\mathbf{r} = \mathbf{m}$: metric shaped by the instantaneous gradient, step along momentum.
- $\mathbf{l} = \mathbf{m}$, $\mathbf{r} = \nabla \mathcal{L}$: metric shaped by accumulated history, step along raw gradient.
- $\mathbf{l} = \mathbf{m}$, $\mathbf{r} = \mathbf{m}$: both metric and step are momentum-based.

7.2 Computational Cost

All variants are $O(N)$ per step (assuming scalar or diagonal γ) and never materialise a full $N \times N$ matrix:

Variant	Dominant operation	Additional metric state
Fixed diagonal	scalar $\times$ vector	none
Learnable scalar	scalar $\times$ vector	1 float (s)
Learnable diagonal	element-wise vector $\times$ vector	N floats (s)
Off-diagonal	inner products + 2×2 inverse	none

The metric-learning overhead for the learnable variants is negligible compared to the forward and backward passes.

7.3 Shared Hyperparameters

The following hyperparameters are common to the fixed diagonal, learnable scalar, and learnable diagonal variants:

- η (**lr**): base learning rate.
- β_m (**momentum**): EMA decay for the momentum buffer. (The paper [1] uses μ for this quantity; we write β_m to avoid collision with the metric learning rate μ .)
- ξ (**xi**): strength of the induced metric correction.
- β (**beta**): EMA decay for the denominator estimate $\hat{v}$.
- λ (**weight_decay**): decoupled weight decay ($\lambda \leq 0$ for decay, following the convention of [1]).

The off-diagonal variant uses η , β_m , ξ , and λ but does not use the denominator EMA β . The learnable variants additionally require:

- μ (**metric_lr**): step size for the online metric update.
- λ_s (**metric_reg**): L2 regularisation on s or s .
- $s_{\max}$ (**metric_clip**): log-domain clipping bound on s , constraining $\alpha = e^s \in [e^{-s_{\max}}, e^{s_{\max}}]$.

Appendix A Repository Structure

A.1 Directory Layout

Directory	Contents
optimisers/	JAX and PyTorch optimiser implementations and equivalence checks
parameters/	Hyperparameter sweep scripts and shared training utilities
analysis/	Jupyter notebooks for loading and visualising sweep results
documents/	This document and supplementary notebooks
images/	Figures used in [1] and this document
results/	Local sweep results (JSON, gitignored)
data/	Local data files (gitignored)

A.2 Running Sweeps

Hyperparameter sweeps are driven by scripts in `parameters/` and support two backends: **local** (Optuna + JSON files, no account required) and **WandB** (cloud logging).

Script	Task	Model
<code>sweep_mnist_mlp.py</code>	MNIST classification	MLP (64 hidden)
<code>sweep_cifar10_resnet18.py</code>	CIFAR-10 classification	ResNet-18
<code>sweep_regression.py</code>	Synthetic regression	MLP
<code>sweep_shake.py</code>	Shakespeare language model	MiniGPT
<code>sweep_small_examples.py</code>	2D test functions	n/a

All scripts share a common CLI:

- `--optimiser`: optimiser name (see below).
- `--num_runs`: number of Optuna trials (default: 50).
- `--backend`: local or wandb (default: wandb).
- `--index`: run index for organising multiple runs of the same sweep.

Results are written to `results/<task>/<optimiser>/run_<index>/<run_id>.json`, each containing the full training history, final summary metrics, and the hyperparameter config used.

A.3 Available Optimisers

Sweep name	Description
<code>adam</code> , <code>adamw</code> , <code>sgd</code> , <code>muon</code>	Baselines
<code>sgd_metric</code> , <code>sgd_log_metric</code> , <code>sgd_rms</code>	Fixed diagonal metric
<code>sgd_learn_scalar</code> , <code>sgd_learn_scalar_log</code>	Learnable scalar metric
<code>sgd_learn_diag</code> , <code>sgd_learn_diag_log</code>	Learnable diagonal metric
<code>sgd_offdiag_{a}-{b}</code>	Off-diagonal metric ($\{0, l, m, \theta\}^2$ pairs, excluding 0.0)

The full optimiser registry is `parameters/optimizer_registry.py`.

A.4 Analysis Notebooks

Notebook	Task
<code>analysis/small_examples.ipynb</code>	2D test functions (Beale, Rosenbrock, etc.)
<code>analysis/mnist_analysis.ipynb</code>	MNIST MLP
<code>analysis/cifar_analysis.ipynb</code>	CIFAR-10 ResNet-18
<code>analysis/regression_analysis.ipynb</code>	Synthetic regression
<code>analysis/shake_analysis.ipynb</code>	Shakespeare language modelling

Each notebook supports both backends. Configure `BACKEND`, `TASK_TAG`, and `OPTIMIZERS` in the second cell. Outputs include training curves, 2D histograms of best metric vs. epoch, time-to-best analysis, and best-hyperparameter tables (also saved as CSV to `analysis/`).

8 Geodesic Convexity Analysis

A natural question arises: does the induced metric $g_{ij} = \gamma_{ij} + \xi l_i l_j$ make the loss landscape “more convex” in a geometric sense? If so, the optimiser’s convergence could be understood through the lens of Riemannian optimisation theory, and stronger guarantees might follow.

The answer turns out to be clean and somewhat surprising. In this section we derive the *Riemannian Hessian* of $\mathcal{L}$ with respect to g in closed form and show that it is a positive scalar multiple of the ordinary Hessian. This means the metric preserves convexity rather than creating it, but the scalar factor provides *adaptive curvature damping* — exactly the smooth gradient clipping behaviour observed empirically.

All formulas below have been verified three ways: (i) abstractly for general N using the xAct tensor algebra system, (ii) by explicit symbolic computation in $n = 2$ and $n = 3$, and (iii) numerically on the benchmark test functions. The full verification notebook is `documents/geodesic-convexity-analysis.nb`.

8.1 Background: Geodesic Convexity

Recall that in standard optimisation a function $\mathcal{L}$ is convex if its Hessian $H_{ij} = \partial_i \partial_j \mathcal{L}$ is positive semi-definite everywhere. The Riemannian generalisation replaces the Euclidean Hessian with the *Riemannian Hessian* (also called the covariant Hessian):

$$(\text{Hess}_g \mathcal{L})_{ij} = \nabla_i \nabla_j \mathcal{L} = \partial_i \partial_j \mathcal{L} - \Gamma_{ij}^k \partial_k \mathcal{L} = H_{ij} - \Gamma_{ij}^k l_k. \quad (19)$$

Here Γ_{ij}^k are the **Christoffel symbols** of g , which encode how the metric’s coordinate basis twists as one moves through parameter space. For a flat (constant) metric the Christoffel symbols vanish and the Riemannian Hessian reduces to the ordinary Hessian. For a curved metric they introduce a *correction term* $\Gamma_{ij}^k l_k$ that subtracts curvature effects from H .

Geodesic Convexity

A function $\mathcal{L}$ is *geodesically convex* with respect to a Riemannian metric g if $\text{Hess}_g \mathcal{L} \succeq 0$ everywhere. If $\mathcal{L}$ is geodesically convex w.r.t. g , then Riemannian gradient descent $\delta \theta^i = -\eta g^{ij} l_j$ is guaranteed to converge to a global minimum — even if $\mathcal{L}$ is non-convex in the ordinary Euclidean sense.

The key question is therefore: can choosing $g = \gamma + \xi l \otimes l$ make the correction term $\Gamma_{ij}^k l_k$ large enough to cancel negative eigenvalues of H and produce a PSD Riemannian Hessian?

8.2 Christoffel Symbols

The Christoffel symbols of the Levi-Civita connection for a metric g are defined by

$$\Gamma_{ij}^k = \frac{1}{2} g^{kl} (\partial_i g_{jl} + \partial_j g_{il} - \partial_l g_{ij}). \quad (20)$$

We compute this in three steps: differentiate the metric, simplify the bracket, and contract with the inverse metric.

Step 1: Metric derivatives. Since γ is constant (parameter-independent), only the rank-one term $\xi l_i l_j$ contributes. Applying the product rule and using $\partial_k l_i = \partial_k \partial_i \mathcal{L} = H_{ki}$:

$$\partial_k g_{ij} = \xi \partial_k (l_i l_j) = \xi (H_{ki} l_j + l_i H_{kj}). \quad (21)$$

Step 2: The Christoffel bracket. The combination inside the parentheses in (20) is called the *Christoffel bracket*. Substituting (21) and writing all six terms explicitly:

$$\begin{aligned} & \partial_i g_{jl} + \partial_j g_{il} - \partial_l g_{ij} \\ &= \xi \left[\underbrace{H_{ij} l_l + l_j H_{il}}_{\partial_i g_{jl}} + \underbrace{H_{ji} l_l + l_i H_{jl}}_{\partial_j g_{il}} - \underbrace{H_{li} l_j + l_i H_{lj}}_{\partial_l g_{ij}} \right]. \end{aligned} \quad (22)$$

Now use the symmetry of the Hessian ($H_{ij} = H_{ji}$, $H_{il} = H_{li}$, $H_{jl} = H_{lj}$). The cross-terms cancel pairwise: $l_j H_{il} - H_{li} l_j = 0$ and $l_i H_{jl} - l_i H_{lj} = 0$. What remains is:

$$\boxed{\partial_i g_{jl} + \partial_j g_{il} - \partial_l g_{ij} = 2 \xi H_{ij} l_l.} \quad (23)$$

This is the key simplification: the bracket is proportional to H_{ij} times l_l , with all information about different index arrangements absorbed by the Hessian symmetry.

Step 3: Contraction with the inverse metric. We need $g^{kl} l_l$, the gradient raised with g (not γ). Starting from the Sherman–Morrison inverse $g^{kl} = \gamma^{kl} - \xi l^k l^l / (1 + \xi \|l\|_\gamma^2)$:

$$\begin{aligned} g^{kl} l_l &= \gamma^{kl} l_l - \frac{\xi l^k l^l l_l}{1 + \xi \|l\|_\gamma^2} = l^k - \frac{\xi l^k \|l\|_\gamma^2}{1 + \xi \|l\|_\gamma^2} = l^k \left(1 - \frac{\xi \|l\|_\gamma^2}{1 + \xi \|l\|_\gamma^2} \right) \\ &= \frac{l^k}{1 + \xi \|l\|_\gamma^2}. \end{aligned} \quad (24)$$

Combining Steps 1–3: $\Gamma_{ij}^k = \frac{1}{2} g^{kl} \cdot 2 \xi H_{ij} l_l = \xi H_{ij} (g^{kl} l_l)$, giving

$$\boxed{\Gamma_{ij}^k = \frac{\xi H_{ij} l^k}{1 + \xi \|l\|_\gamma^2}} \quad (25)$$

where $l^k = \gamma^{kk} l_k$ is the γ -raised gradient and $\|l\|_\gamma^2 = \gamma^{mn} l_m l_n$.

Sanity checks.

- **Symmetry:** $\Gamma_{ij}^k = \Gamma_{ji}^k$ since $H_{ij} = H_{ji}$. ✓
- **Flat limit:** when $\xi = 0$, $\Gamma_{ij}^k = 0$, recovering the flat connection of γ . ✓
- **Proportionality:** Γ_{ij}^k is proportional to H_{ij} for all k . This will drive the main result.

8.3 Riemannian Hessian

The correction term. Substituting (25) into the correction $\Gamma_{ij}^k l_k$:

$$\Gamma_{ij}^k l_k = \frac{\xi H_{ij} l^k l_k}{1 + \xi \|l\|_\gamma^2} = \frac{\xi \|l\|_\gamma^2}{1 + \xi \|l\|_\gamma^2} H_{ij}. \quad (26)$$

The correction is proportional to H_{ij} with scalar coefficient $\xi \|l\|_\gamma^2 / (1 + \xi \|l\|_\gamma^2) \in [0, 1)$.

Main result. The Riemannian Hessian (19) becomes

$$(\text{Hess}_g \mathcal{L})_{ij} = H_{ij} - \frac{\xi \|l\|_\gamma^2}{1 + \xi \|l\|_\gamma^2} H_{ij} = H_{ij} \left(1 - \frac{\xi \|l\|_\gamma^2}{1 + \xi \|l\|_\gamma^2} \right), \quad (27)$$

and since $1 - \frac{\xi \|l\|_\gamma^2}{1 + \xi \|l\|_\gamma^2} = \frac{1}{1 + \xi \|l\|_\gamma^2}$, we obtain

$$\boxed{(\text{Hess}_g \mathcal{L})_{ij} = \frac{H_{ij}}{1 + \xi \|l\|_\gamma^2}} \quad (28)$$

Interpretation

The Riemannian Hessian is the Euclidean Hessian divided by the positive scalar $(1 + \xi \|l\|_\gamma^2)$. This is the same denominator that appears in the optimiser's update rule (4). The metric does not rotate eigenvectors or mix curvature directions — it applies a uniform, position-dependent damping.

Eigenvalue analysis. Let $\lambda_1, \dots, \lambda_N$ be the eigenvalues of H . Since $\text{Hess}_g \mathcal{L} = H/(1 + \xi \|l\|_\gamma^2)$ is a scalar multiple, its eigenvalues are $\lambda_k/(1 + \xi \|l\|_\gamma^2)$ with the same eigenvectors. The scaling factor $1/(1 + \xi \|l\|_\gamma^2) \in (0, 1]$ is strictly positive, so:

1. **Sign preservation.** Every eigenvalue keeps its sign: positive stays positive, negative stays negative, zero stays zero. *The induced metric cannot flip a saddle direction into a convex one.*
2. **Geodesic convexity $\iff$ Euclidean convexity.** $\text{Hess}_g \mathcal{L} \succeq 0$ if and only if $H \succeq 0$. The rank-one induced metric cannot convert a non-convex function into a geodesically convex one, regardless of ξ .
3. **Condition number is unchanged.** $\kappa(\text{Hess}_g \mathcal{L}) = \lambda_{\max}/\lambda_{\min} = \kappa(H)$: the uniform scaling cancels in the ratio.
4. **Magnitude damping.** Although signs are preserved, *all curvatures are damped* by a factor that depends on position:
 - Far from critical points ($\|l\|$ large): damping is strong, $|\lambda_k| \rightarrow 0$.
 - Near critical points ($\|l\| \rightarrow 0$): damping vanishes, $\text{Hess}_g \mathcal{L} \rightarrow H$.

The metric flattens the landscape in steep regions while preserving the full curvature near minima and saddle points.

Worked example: quadratic loss. For $\mathcal{L}(\theta) = \frac{1}{2} \theta^\top A \theta$ with symmetric A , we have $H = A$ (constant), $l = A\theta$, and $\|l\|^2 = \theta^\top A^\top A \theta$. With $\gamma = \mathbb{1}_N$:

$$(\text{Hess}_g \mathcal{L})_{ij} = \frac{A_{ij}}{1 + \xi \|A\theta\|^2}. \quad (29)$$

Near the minimum $\theta = 0$: $\|A\theta\| \rightarrow 0$, so $\text{Hess}_g \mathcal{L} \rightarrow A$. Far from the minimum: $\text{Hess}_g \mathcal{L} \rightarrow 0$, regardless of whether A is PSD or indefinite. If A has a negative eigenvalue (saddle), that eigenvalue remains negative everywhere but is progressively damped away from the origin.

8.4 Connection to the Gauss Equation

The same results can be obtained from a complementary perspective: the *extrinsic geometry* of the loss surface. This route gives additional insight and connects the Hessian to curvature in a way that is geometrically transparent.

Setup. The embedding $\phi(\theta) = (\theta^1, \dots, \theta^N, \mathcal{L}(\theta))$ defines a hypersurface in $\mathbb{R}^{N+1}$. The ambient metric is $h = \text{diag}(\gamma_1, \dots, \gamma_N, \xi)$. The tangent vectors are $e_i = (\delta_{i1}, \dots, \delta_{iN}, l_i)$, and the induced metric is $g_{ij} = h(e_i, e_j) = \gamma_{ij} + \xi l_i l_j$ as expected.

Unit normal. The unit normal n (with respect to h) satisfies $h(n, e_i) = 0$ and $h(n, n) = 1$. Solving these constraints gives

$$n = \frac{1}{\sqrt{\xi(1 + \xi \|l\|_\gamma^2)}} \left(-\xi l^1, \dots, -\xi l^N, 1 \right). \quad (30)$$

We write $n_{N+1} := 1/\sqrt{\xi(1 + \xi \|l\|_\gamma^2)}$ for the component along the loss axis.

Second fundamental form. The second fundamental form measures how the surface curves within the ambient space. For our embedding, $\partial_i \partial_j \phi = (0, \dots, 0, H_{ij})$, so

$$b_{ij} = h(\partial_i \partial_j \phi, n) = \xi n_{N+1} H_{ij} = \sqrt{\frac{\xi}{1 + \xi \|l\|_\gamma^2}} H_{ij}. \quad (31)$$

Like the Riemannian Hessian, b_{ij} is proportional to H_{ij} . The two are related by

$$(\text{Hess}_g \mathcal{L})_{ij} = n_{N+1} b_{ij}, \quad (32)$$

connecting the covariant Hessian to the extrinsic curvature of the loss surface.

Gauss equation. For a hypersurface in a flat ambient space, the Gauss equation expresses the intrinsic Riemann curvature tensor entirely in terms of the second fundamental form:

$$R_{ijkl} = b_{ik} b_{jl} - b_{il} b_{jk} = \frac{\xi}{1 + \xi \|l\|_\gamma^2} (H_{ik} H_{jl} - H_{il} H_{jk}). \quad (33)$$

This tensor inherits all the standard symmetries (verified abstractly using xAct): skew-symmetry in each index pair, pair symmetry $R_{ijkl} = R_{klij}$, and the first Bianchi identity $R_{ijkl} + R_{iklj} + R_{iljk} = 0$.

Gaussian curvature. For a 2D parameter space ($N = 2$), the only independent component of R_{ijkl} is R_{1212} , and the Gaussian curvature is

$$K = \frac{R_{1212}}{\det(g)} = \frac{\xi \det(H)}{\det(\gamma) (1 + \xi \|l\|_\gamma^2)^2}. \quad (34)$$

The Ricci tensor satisfies $R_{ij} = K g_{ij}$ and the scalar curvature is $R = 2K$, as expected for a surface.

Curvature–Hessian dictionary. Equation (34) provides a direct link between the intrinsic geometry of the parameter surface and the loss Hessian:

Surface geometry	Condition	Hessian meaning
Positive curvature ($K > 0$)	$\det(H) > 0$	H definite (bowl or dome)
Negative curvature ($K < 0$)	$\det(H) < 0$	H indefinite (saddle)
Zero curvature ($K = 0$)	$\det(H) = 0$	H singular (ridge/valley)

The magnitude of K is controlled by ξ : larger ξ amplifies curvature near critical points ($\|l\|$ small) and suppresses it in steep regions ($\|l\|$ large).

8.5 Off-Diagonal Metric

The off-diagonal metric of Section 6 has the form $g_{ij} = \gamma_{ij} + \xi(a_i l_j + b_j l_i + l_i l_j)$ with coupling vectors a, b . A natural question is whether this richer structure can achieve what the rank-one metric cannot.

Bracket derivation. With a, b treated as constant vectors, $\partial_k g_{ij} = \xi[(a_i + l_i)H_{kj} + (b_j + l_j)H_{ki}]$. Writing $u_i := a_i + l_i$ and $v_j := b_j + l_j$ for brevity, the six bracket terms become:

$$\begin{aligned} & \partial_i g_{jl} + \partial_j g_{il} - \partial_l g_{ij} \\ &= \xi[u_j H_{il} + v_l H_{ij} + u_i H_{jl} + v_l H_{ij} - u_i H_{lj} - v_j H_{li}] \\ &= \xi[(u_j - v_j)H_{il} + 2v_l H_{ij}] \\ &= \xi[(a_j - b_j)H_{il} + 2(b_l + l_l)H_{ij}]. \end{aligned} \tag{35}$$

Here we used $H_{jl} = H_{lj}$ and $H_{il} = H_{li}$ to cancel terms, and substituted back $u - v = a - b$ and $v_l = b_l + l_l$.

Two cases.

1. **Symmetric case** ($a = b$). The $(a_j - b_j)$ term vanishes, leaving $2\xi(a_l + l_l)H_{ij}$ — again proportional to H_{ij} . The Riemannian Hessian is a positive scalar multiple of H (with a different scalar from the rank-one case), and the same sign-preservation result holds.
2. **Asymmetric case** ($a \neq b$). The term $(a_j - b_j)H_{il}$ mixes Hessian components in different index directions. The Riemannian Hessian correction is no longer proportional to H , and its eigenvalues can differ qualitatively from those of H . A numerical scan over $\sim 14,000$ configurations of (a, b, l, ξ) with an indefinite H found:

- **58 cases** where the Riemannian Hessian became PSD (saddle $\rightarrow$ geodesically convex).
- **But:** in all 58 cases the metric g was *not symmetric* ($g_{ij} \neq g_{ji}$), so g is not a Riemannian metric.

The symmetry condition $g_{ij} = g_{ji}$ requires $a_i l_j + l_i b_j = a_j l_i + l_j b_i$ for all i, j , which holds if and only if $a = b$ (or $a - b \propto l$). A targeted scan over all symmetric off-diagonal configurations found **zero cases** where a saddle became convex.

8.6 Learnable Metrics: Genuine Curvature Correction

The results above establish that the fixed metric (γ constant) can only rescale the Hessian, not correct its eigenvalue structure. We now show that this limitation is *specific to constant* γ : when γ depends on θ (Sections 4–5), the Riemannian Hessian acquires additional terms that can change eigenvalue signs.

General formula. When $\gamma_{ij}(\theta)$ varies with the parameters, the metric derivative gains an extra contribution:

$$\partial_k g_{ij} = \partial_k \gamma_{ij} + \xi(H_{ki} l_j + l_i H_{kj}). \tag{36}$$

The Christoffel bracket becomes:

$$\partial_i g_{jl} + \partial_j g_{il} - \partial_l g_{ij} = \underbrace{[\partial_i \gamma_{jl} + \partial_j \gamma_{il} - \partial_l \gamma_{ij}]}_{\text{bracket of } \gamma} + \underbrace{2\xi H_{ij} l_l}_{\text{rank-one term}}. \tag{37}$$

The rank-one term is exactly as before; the new ingredient is the Christoffel bracket of γ itself. Contracting with $\frac{1}{2}g^{kl}$ and then with l_k , the Riemannian Hessian becomes (after simplification; see the verification notebook for the full derivation):

$$\boxed{(\text{Hess}_g \mathcal{L})_{ij} = \frac{H_{ij} - C_{ij}}{1 + \xi \|l\|_\gamma^2}} \tag{38}$$

where the **curvature correction** is

$$C_{ij} := \Gamma_{ij}^k(\gamma) l_k. \quad (39)$$

For constant γ , $C_{ij} = 0$ and we recover (28). For θ -dependent γ , C_{ij} depends on the derivatives of γ and on the gradient l — but *not* on H . This means the correction can have an eigenvalue structure completely different from that of H .

Learnable scalar: $\gamma = e^s \mathbb{1}_N$. For $\gamma_{ij} = e^{s(\theta)} \delta_{ij}$ with $s_k := \partial_k s$, the Christoffel symbols of γ are those of a conformally flat metric:

$$\Gamma_{ij}^k(\gamma) = \frac{1}{2} (s_i \delta_{jk} + s_j \delta_{ik} - s_k \delta_{ij}). \quad (40)$$

The correction term is:

$$C_{ij} = \Gamma_{ij}^k(\gamma) l_k = \frac{1}{2} (s_i l_j + s_j l_i - (\nabla s \cdot l) \delta_{ij}). \quad (41)$$

This is a symmetric matrix built from the outer products of ∇s and l . Its trace is

$$\text{tr}(C) = \delta^{ij} C_{ij} = \nabla s \cdot l - \frac{N}{2} (\nabla s \cdot l) = \frac{2 - N}{2} (\nabla s \cdot l). \quad (42)$$

For $N = 2$ the correction is traceless, and for $N > 2$ the trace is *negative* (when $\nabla s \cdot l > 0$), meaning the correction shifts the *average* eigenvalue of the Riemannian Hessian downward.

The eigenvalues of C scale as $\pm \|\nabla s\| \|l\|/2$ (the precise values depend on the angle between ∇s and l). For $N = 2$ the tracelessness forces C to redistribute curvature between directions without shifting the mean, so a sign flip requires $\text{tr}(H) > 0$. For $N > 2$ the situation is *harder*: the negative trace pulls eigenvalues down on average, so even larger $\text{tr}(H)$ is needed to compensate. Numerical tests confirm that the required $\text{tr}(H)$ for a sign flip grows with N , strengthening the hierarchy.

Learnable diagonal: $\gamma = \text{diag}(e^{s_1}, \dots, e^{s_N})$. For $\gamma_{ij} = e^{s_i} \delta_{ij}$ (no sum) with $s_{i,k} := \partial_k s_i$, the Christoffel symbols of γ are:

$$\Gamma_{ij}^k(\gamma) = \frac{1}{2} (s_{j,i} \delta_{jk} + s_{i,j} \delta_{ik} - e^{s_i - s_k} s_{i,k} \delta_{ij}), \quad (43)$$

where the exponential factor $e^{s_i - s_k}$ arises from raising an index with $\gamma^{kk} = e^{-s_k}$. The correction $C_{ij} = \Gamma_{ij}^k(\gamma) l_k$ now has **non-zero trace**, with contributions amplified by the exponential ratios $e^{s_i - s_k}$.

This is strictly more powerful than the scalar case:

- The non-zero trace means C can shift *all* eigenvalues in the same direction, not just redistribute.
- The exponential factors allow large corrections when the per-parameter scales s_i differ significantly.
- Sign flips are possible even when $\text{tr}(H) \leq 0$ (predominantly concave loss).

Numerical verification. On the saddle function $\mathcal{L} = x^2 - y^2$ (with $H = \text{diag}(2, -2)$, $\text{tr}(H) = 0$):

Metric	eig(H)	tr(C)	eig(Hess $_g \mathcal{L}$)	Sign flip?
Constant γ	$\{-2, 2\}$	0	$\{-0.05, 0.05\}$	No
Scalar $e^s \mathbb{1}$	$\{-2, 2\}$	0	$\{0, 0\}$	No (degenerate)
Diagonal $\text{diag}(e^{s_i})$	$\{-2, 2\}$	-4	$\{+0.08, +0.25\}$	Yes

With an appropriately chosen ∇s , the diagonal learnable metric converts an indefinite Hessian into a positive-definite Riemannian Hessian. This was also verified on more extreme cases: $\mathcal{L} = x^2 - 3y^2$ ($\text{tr}(H) = -4$) and $\mathcal{L} = x^2 - 100y^2$ ($\text{tr}(H) = -198$), both achieving sign flips.

Caveat: theoretical capability vs. current optimizer. The sign flips above use ∇s values chosen to maximize $\lambda_{\min}$. The *current* online learning rule $s_i \leftarrow s_i + \mu \xi e^{s_i} l_i^2 - \mu \lambda_s s_i$ does *not* produce these optimal gradients. At symmetric saddle points (where l_i^2 is equal across directions), the rule drives all s_i toward the same value, creating zero trace correction and no sign flip. Numerical tests confirm that the online rule fails to achieve sign flips at the saddle $\mathcal{L} = x^2 - y^2$ for all tested values of λ_s .

The gap between theoretical capability and current optimizer behavior is substantial: the optimal ∇s achieves $\lambda_{\min} \approx +0.36$ at the saddle where the online rule gives $\lambda_{\min} \approx -0.07$. Closing this gap — for instance by incorporating curvature-sign information into the metric update — is a promising direction for future work.

Key Finding: Hierarchy of Curvature Correction

There is a strict hierarchy of curvature correction power:

Constant γ < Scalar learnable < Diagonal learnable

- *Constant γ* : $C = 0$. Only rescaling, never changes eigenvalue signs.
- *Scalar learnable* ($e^s \mathbb{1}$): $\text{tr}(C) = \frac{2-N}{2}(\nabla s \cdot l)$. For $N = 2$: traceless, can flip signs when $\Delta \mathcal{L} > 0$. For $N > 2$: negative trace makes flipping *harder*; the required $\text{tr}(H)$ grows with N .
- *Diagonal learnable* ($\text{diag}(e^{s_i})$): C has non-zero trace amplified by $e^{s_i - s_k}$. Can flip signs even for predominantly concave losses, independent of N .

The learnable diagonal metric performs **genuine Riemannian curvature correction**, not just gradient clipping. This provides a geometric explanation for its empirical superiority (Section 5).

Practical implications. The metric learning rule of Section 5 updates s_i via gradient ascent on the denominator trace: $s_i \leftarrow s_i + \mu \xi e^{s_i} l_i^2 - \mu \lambda_s s_i$. This rule was designed to maximise adaptive step-size calibration, not to optimise the curvature correction C_{ij} . The analysis above shows that ∇s produced by this rule is suboptimal for sign flipping: the optimal ∇s requires anti-correlation with H_{ii} (shrink scale along positive-curvature directions, grow along negative-curvature), whereas the online rule drives s_i proportionally to l_i^2 regardless of curvature sign. An update rule that incorporates curvature-sign information — for instance via secant approximations to H_{ii} from consecutive gradient differences — could close the gap between theoretical capability and actual performance.

8.7 Implications for the Optimiser

The analysis above reveals a two-regime picture:

Fixed metric: adaptive step-size rescaling. The central result (28) shows that the fixed induced metric $g = \gamma + \xi l \otimes l$ (with constant γ) provides *adaptive step-size rescaling* rather than curvature correction. The parameter update (1),

$$\delta \theta^i = -\eta g^{ij} l_j = -\eta \frac{l^i}{1 + \xi \|l\|_\gamma^2}, \quad (44)$$

moves along the same direction as preconditioned gradient descent with γ , but with the effective learning rate reduced by $1/(1 + \xi \|l\|_\gamma^2)$. This is the “smooth gradient clipping” interpretation:

- When $\|l\| \gg 1/\sqrt{\xi}$: the step size scales as $\sim 1/(\xi \|l\|)$, preventing large steps in steep regions.
- When $\|l\| \ll 1/\sqrt{\xi}$: the step size approaches the unscaled ηl^i , allowing full steps near critical points where the gradient is informative.

Learnable metric: curvature correction. When γ is learned online (Sections 4–5), the Riemannian Hessian (38) acquires the correction C_{ij} from the curvature of γ itself. The optimiser is no longer just clipping gradients — it is actively reshaping the effective curvature of the loss landscape. The diagonal learnable variant is the most powerful, capable of converting saddle directions into convex ones. This may explain why `sgd_learn_diag` outperforms all other variants in the benchmarks, including Adam.

Open direction: log-loss embedding. The variant $f(\theta) = \log \mathcal{L}(\theta)$ changes the gradient from l_i to $l_i/\mathcal{L}$ and modifies the Hessian structure, potentially yielding further curvature properties. This remains to be analysed.

Figures 6–9 illustrate the eigenvalue scaling and curvature damping on the Rosenbrock and saddle functions.




Figure 6: Minimum eigenvalue of the Euclidean Hessian H and the Riemannian Hessian $\text{Hess}_g \mathcal{L}$ along a semicircular arc through the Rosenbrock landscape, for several values of ξ . The sign is preserved; only the magnitude is damped.

References

- [1] T. R. Harvey, “The Optimiser Hidden in Plain Sight: Training with the Loss Landscape’s Induced Metric,” *arXiv preprint arXiv:2509.03594*, 2025.




Figure 7: The scaling factor $1/(1 + \xi \|\nabla \mathcal{L}\|^2)$ along the Rosenbrock valley $y = x^2$. At the minimum ($t = 0.5$, corresponding to the point $(1, 1)$) where $\|\nabla \mathcal{L}\| \rightarrow 0$, the factor approaches 1; away from the minimum it is strongly suppressed, with larger ξ producing a narrower peak.




Figure 8: Minimum eigenvalue of H (left) and $\text{Hess}_g \mathcal{L}$ with $\xi = 10$ (right) over the Rosenbrock domain $[-2, 2]^2$. The sign structure is identical; the induced metric compresses the magnitude in high-gradient regions.




Figure 9: Eigenvalues of H and $\text{Hess}_g \mathcal{L}$ ($\xi = 10$) for the saddle function $\mathcal{L} = x^2 - y^2$ on the unit circle. The Euclidean eigenvalues ± 2 (dashed) are uniformly compressed but never change sign.