

Induced Metric Optimisers

March 3, 2026

Contents

1	Problem Statement	2
2	Fixed Diagonal Metric	2
3	Learnable Scalar Metric	3
4	Learnable Diagonal Metric	5
5	Off-Diagonal Metric	6
6	Implementation	8
7	Results	11

1 Problem Statement

Given a model with N parameters $\theta(t) = \{\theta^i(t)\}_{i=1}^N$ and a loss function $\mathcal{L} : \mathbb{R}^N \rightarrow \mathbb{R}$, our goal is to find the most optimal parameter update $\delta\theta$ at each time step t . That is, in gradient descent we have

$$\frac{d\theta^i}{dt} = -g^{ij} \frac{\partial \mathcal{L}}{\partial \theta^j} \implies \delta\theta^i \approx -\eta g^{ij} \frac{\partial \mathcal{L}}{\partial \theta^j}$$

where $\eta \in \mathbb{R}_{>0}$ is the learning rate and g^{ij} is the inverse of the induced metric on the parameter space.

Across all variants in this document the induced metric is obtained by the same geometric construction: we embed parameter space into an ambient space via

$$\phi : \mathbb{R}^N \rightarrow \mathbb{R}^{N+1}, \quad \phi(\theta^1, \dots, \theta^N) = (\theta^1, \dots, \theta^N, f(\theta)),$$

where $f : \mathbb{R}^N \rightarrow \mathbb{R}$ is a scalar function of the parameters (taken to be the loss $\mathcal{L}$, or a function thereof), and then pull back an ambient bilinear form $h \in \mathbb{R}^{(N+1) \times (N+1)}$ to the parameter space. The variants differ in the choice of h .

We use the shorthand $l_i := \partial \mathcal{L} / \partial \theta^i$.

2 Fixed Diagonal Metric

The original ambient metric in the paper is

$$h = \begin{pmatrix} \gamma & \vec{0} \\ \vec{0}^\top & 1 \end{pmatrix} \in \mathbb{R}^{(N+1) \times (N+1)}$$

where $\gamma \in \mathbb{R}^{N \times N}$. The induced metric g on the parameter space is given by

$$g_{ij} = h_{\alpha\beta} \frac{\partial \phi^\alpha}{\partial \theta^i} \frac{\partial \phi^\beta}{\partial \theta^j} = \gamma_{ij} + \frac{\partial \mathcal{L}}{\partial \theta^i} \frac{\partial \mathcal{L}}{\partial \theta^j}.$$

Then, to calculate the parameter updates $\delta\theta^i$ we need g^{ij} , the inverse of g_{ij} . We can compute this inverse using the Sherman-Morrison-Woodbury formula:

Sherman-Morrison-Woodbury Formula

Let $A \in \mathbb{R}^{n \times n}$ be an invertible matrix, and let $U, V \in \mathbb{R}^{n \times k}$ be matrices. Then, if $C = A + UV^\top$ is invertible, we have

$$C^{-1} = A^{-1} - A^{-1}U(I_k + V^\top A^{-1}U)^{-1}V^\top A^{-1}.$$

This gives us

$$g^{ij} = \gamma^{ij} - \frac{\gamma^{ik} \gamma^{jl} \frac{\partial \mathcal{L}}{\partial \theta^k} \frac{\partial \mathcal{L}}{\partial \theta^l}}{1 + \gamma^{mn} \frac{\partial \mathcal{L}}{\partial \theta^m} \frac{\partial \mathcal{L}}{\partial \theta^n}} = \gamma^{ij} - \frac{l^i l^j}{1 + l_k l^k}$$

where

$$l_i := \frac{\partial \mathcal{L}}{\partial \theta^i} \implies l^i = \gamma^{ij} l_j.$$

Thus, the parameter updates are given by

$$\begin{aligned} \delta \theta^i &= -\eta \left(\gamma^{ij} - \frac{l^i l^j}{1 + l_k l^k} \right) l_j \\ &= -\eta l^i \left(1 - \frac{l_j l^j}{1 + l_k l^k} \right) \\ &= -\eta l^i \left(\frac{1}{1 + l_k l^k} \right). \end{aligned}$$

So to summarize, (1) we start with an ambient metric h which induces a metric g on the parameter space, (2) we compute the inverse g^{ij} , and (3) we use g^{ij} to compute the parameter updates $\delta \theta^i$.

3 Learnable Scalar Metric

Construction

In the fixed diagonal variant, the base metric γ is a fixed hyperparameter. Here we instead allow γ to be *learned online* during optimisation, starting with the simplest possible form: a single scalar. We parameterise the inverse metric as

$$\gamma_{\phi}^{ij} = \alpha \delta^{ij}, \quad \alpha = e^s,$$

where $s \in \mathbb{R}$ is a single learnable scalar. Initialising $s = 0$ gives $\gamma^{-1} = \mathbb{1}_N$, so the optimiser starts as the unscaled fixed diagonal variant.

The ambient metric remains block-diagonal,

$$h = \begin{pmatrix} e^{-s} \mathbb{1}_N & \vec{0} \\ \vec{0}^\top & 1 \end{pmatrix},$$

and the induced metric on parameter space is

$$g_{ij} = e^{-s} \delta_{ij} + l_i l_j.$$

Applying the Sherman-Morrison-Woodbury formula with $A_{ij} = e^{-s} \delta_{ij}$, $u = v = \mathbf{l}$, we get

$$g^{ij} = e^s \delta^{ij} - \frac{e^{2s} l_i l_j}{1 + e^s \|\mathbf{l}\|^2}.$$

The parameter update is therefore

$$\delta \theta^i = -\eta g^{ij} l_j = -\eta e^s l_i \left(1 - \frac{e^s \|\mathbf{l}\|^2}{1 + e^s \|\mathbf{l}\|^2} \right) = -\frac{\eta e^s l_i}{1 + e^s \|\mathbf{l}\|^2}.$$

Adding the hyperparameter ξ to control the overall strength of the metric correction, and using an exponential moving average (EMA) of the denominator quantity for stability (as in Algorithms 1 and 2 of the paper), the practical update is

$$\delta\theta^i = -\eta r_t (\alpha m^i), \quad r_t = \frac{1}{1 + |\hat{v}_t|}, \quad v_t = \xi \alpha \|\mathbf{g}_t\|^2, \quad (1)$$

where $\mathbf{m}$ is the bias-corrected momentum vector and $\hat{v}_t$ is the bias-corrected EMA of v_t . For the log-loss embedding (Algorithm 2) the damping factor becomes $r_t = \mathcal{L}_t / (\mathcal{L}_t^2 + |\hat{v}_t|)$.

Metric Learning

We want to learn s so that $\alpha = e^s$ is calibrated to the local geometry of the loss. The trace quantity $v = \xi \alpha \|\mathbf{g}\|^2$ already captures how strongly the loss surface varies along the gradient direction under the current metric. We treat v as a surrogate objective and perform a *gradient ascent* step on s :

$$\frac{\partial v}{\partial s} = \xi e^s \|\mathbf{g}\|^2 = v.$$

Thus the online metric update is

$$s \leftarrow s + \mu v - \mu \lambda s \quad (2)$$

where $\mu > 0$ is the metric learning rate and $\lambda > 0$ is an L2 regularisation coefficient. The regularisation term $-\mu \lambda s$ penalises large deviations of s from zero (i.e. it pulls the inverse metric back towards the identity) and ensures stability. Finally s is clipped to the interval $[-s_{\max}, s_{\max}]$.

Intuitively: if the gradient is large, v is large, so s increases and α grows. A larger α amplifies the preconditioned direction $\alpha \mathbf{m}$ in the numerator, but also increases v (and hence $\hat{v}$) which shrinks r_t in the denominator. The net effect is that the metric adapts the scale of updates dynamically, with the regulariser providing a restoring force towards the identity.

Hyperparameters

The learnable scalar variant introduces three new hyperparameters beyond those of the fixed diagonal variant:

- μ (**metric_lr**): step size for the online metric update in (2).
- λ (**metric_reg**): L2 regularisation strength on s , controlling how far the metric drifts from identity.
- $s_{\max}$ (**metric_clip**): clipping bound on s , bounding $\alpha \in [e^{-s_{\max}}, e^{s_{\max}}]$.

4 Learnable Diagonal Metric

Construction

The learnable scalar metric uses a single degree of freedom and therefore cannot differentiate between parameters. We now promote s to a full vector $\mathbf{s} \in \mathbb{R}^N$ with one entry per parameter, giving the per-parameter inverse metric

$$\gamma_{\phi}^{ij} = e^{s_i} \delta^{ij}.$$

The ambient metric is

$$h = \begin{pmatrix} \text{diag}(e^{-\mathbf{s}}) & \vec{0} \\ \vec{0}^\top & 1 \end{pmatrix},$$

and the induced metric on parameter space is

$$g_{ij} = e^{-s_i} \delta_{ij} + l_i l_j.$$

This is again a rank-1 update to a diagonal matrix. Applying Sherman-Morrison-Woodbury with $A_{ij} = e^{-s_i} \delta_{ij}$ and $u = v = \mathbf{l}$,

$$g^{ij} = e^{s_i} \delta^{ij} - \frac{e^{s_i} l_i e^{s_j} l_j}{1 + \sum_k e^{s_k} l_k^2}.$$

The update is

$$\delta\theta^i = -\eta g^{ij} l_j = -\eta \left(e^{s_i} l_i - \frac{e^{s_i} l_i \sum_k e^{s_k} l_k^2}{1 + \sum_k e^{s_k} l_k^2} \right) = -\frac{\eta e^{s_i} l_i}{1 + \sum_k e^{s_k} l_k^2}.$$

With ξ , momentum $\mathbf{m}$, and the EMA damping, the practical update is

$$\delta\theta^i = -\eta r_t (e^{s_i} m^i), \quad r_t = \frac{1}{1 + |\hat{v}_t|}, \quad v_t = \xi \sum_k e^{s_k} g_k^2, \quad (3)$$

with the same log-loss variant as before. Compared to the scalar case, each parameter θ^i now has an individual scale e^{s_i} , allowing the metric to adapt separately to different regions of parameter space.

Metric Learning

The per-parameter metric update treats each s_i independently. The surrogate objective is

$$J(\mathbf{s}) = \xi \sum_k e^{s_k} g_k^2,$$

and its gradient with respect to s_i is

$$\frac{\partial J}{\partial s_i} = \xi e^{s_i} g_i^2.$$

The online update is therefore

$$s_i \leftarrow s_i + \mu \xi e^{s_i} g_i^2 - \mu \lambda s_i \quad (4)$$

followed by mean-centering and clipping (discussed below).

Mean-Centering

After the gradient step (4), we apply a per-leaf mean-centering operation:

$$s_i \leftarrow s_i - \bar{s}, \quad \bar{s} = \frac{1}{N} \sum_k s_k.$$

This is not an ad-hoc trick but a principled normalisation. The trace $v = \xi \sum_k e^{s_k} g_k^2$ has a global scale degeneracy: shifting all s_k by a constant c is equivalent to scaling ξ by e^c and leaving $\mathbf{s}$ unchanged. Concretely, writing $s_k = \bar{s} + \delta_k$ (so $\sum_k \delta_k = 0$),

$$v = \xi e^{\bar{s}} \sum_k e^{\delta_k} g_k^2.$$

The global mean $\bar{s}$ is exactly degenerate with ξ . Mean-centering enforces $\bar{s} = 0$ after every step, which fixes this gauge and ensures that ξ remains the single scalar controlling the overall metric magnitude while $\mathbf{s}$ captures only the *relative* per-parameter geometry.

Scalar as a Special Case

If all s_i are initialised to the same value and the gradients have the same magnitude for all parameters at every step, then $\mathbf{s}$ stays uniform and mean-centering keeps it at zero. In that case $e^{s_i} = 1$ for all i and $v = \xi \|\mathbf{g}\|^2$, which reduces exactly to the fixed diagonal (scalar $\gamma = 1$) variant. The learnable diagonal metric therefore strictly generalises the fixed and scalar cases.

5 Off-Diagonal Metric

Now, we consider a more general ambient bilinear form of the form

$$h = \begin{pmatrix} \gamma & \vec{a} \\ \vec{b}^\top & 1 \end{pmatrix} \in \mathbb{R}^{(N+1) \times (N+1)}$$

where $\vec{a}, \vec{b} \in \mathbb{R}^N$. Intuitively, the vectors $\vec{a}$ and $\vec{b}$ introduce off-diagonal couplings between the parameter space and the loss function dimension. In general, if $\vec{a} \neq \vec{b}$ this h is not symmetric, so here we simply treat it as a preconditioning matrix rather than a Riemannian metric. The induced bilinear form on the parameter space is then given by

$$\begin{aligned} g_{ij} &= h_{\alpha\beta} \frac{\partial \phi^\alpha}{\partial \theta^i} \frac{\partial \phi^\beta}{\partial \theta^j} \\ &= \gamma_{ij} + a_i \frac{\partial \mathcal{L}}{\partial \theta^j} + b_j \frac{\partial \mathcal{L}}{\partial \theta^i} + \frac{\partial \mathcal{L}}{\partial \theta^i} \frac{\partial \mathcal{L}}{\partial \theta^j} \\ &= \gamma_{ij} + a_i l_j + b_j l_i + l_i l_j. \end{aligned}$$

To compute the inverse g^{ij} , we notice that g can be expressed as a rank-2 update to γ , that is for $U = [a, l]$, $V = [l, b + l] \in \mathbb{R}^{N \times 2}$ we have

$$\begin{aligned}\gamma_{ij} + U_{ik}V_{jk} &= \gamma_{ij} + a_i l_j + (b_j + l_j) l_i \\ &= \gamma_{ij} + a_i l_j + b_j l_i + l_i l_j \\ &= g_{ij}.\end{aligned}$$

Thus, we can apply the Sherman-Morrison-Woodbury formula again to get

$$\begin{aligned}g^{ij} &= \gamma^{ij} - \gamma^{ik} U_{k\alpha} \left(\delta_{\alpha\beta} + V_{m\alpha} \gamma^{mn} U_{n\beta} \right)^{-1} V_{l\beta} \gamma^{lj} \\ &= \gamma^{ij} - \gamma^{ik} (a_k, l_k) \begin{pmatrix} 1 + l_n a^n & l_n l^n \\ (b_m + l_m) \gamma^{mn} a_n & 1 + (b_m + l_m) \gamma^{mn} l_n \end{pmatrix}^{-1} \begin{pmatrix} l_n \\ b_n + l_n \end{pmatrix} \gamma^{nj}.\end{aligned}$$

For convenience we define the following scalar contractions:

$$c_{al} := a_n l^n, \quad c_{ll} := l_n l^n, \quad c_{ba} := b_n a^n, \quad c_{bl} := b_n l^n.$$

Thus, we have

$$\begin{aligned}g^{ij} &= \gamma^{ij} - \frac{\gamma^{ik}}{D} \left[a_k ((1 + c_{bl}) l_n - c_{ll} b_n) \right. \\ &\quad \left. + l_k (-(c_{ba} + c_{al}) l_n + (1 + c_{al})(b_n + l_n)) \right] \gamma^{nj}\end{aligned}$$

where

$$\begin{aligned}D &= (1 + c_{al})(1 + c_{bl} + c_{ll}) - c_{ll}(c_{ba} + c_{al}) \\ &= 1 + c_{al} + c_{bl} + c_{al}c_{bl} + c_{ll} - c_{ll}c_{ba}.\end{aligned}$$

A numerical check of this has been performed in documents/off-diag-check.ipynb.

Quick Sanity Check

If we set $\vec{a} = \vec{0}$ and $\vec{b} = \vec{0}$, we should recover the diagonal metric case. In this case, $c_{al} = 0$, $c_{ba} = 0$, $c_{bl} = 0$, and hence $D = 1 + c_{ll}$. Then

$$g^{ij} = \gamma^{ij} - \gamma^{ik} \frac{1}{1 + c_{ll}} l_k l_n \gamma^{nj} = \gamma^{ij} - \frac{l^i l^j}{1 + l_k l^k},$$

which matches our previous result.

The parameter updates are given by

$$\delta\theta^i = -\eta g^{ij} l_j = -\eta \left[\frac{1 + c_{al}}{D} l^i - \frac{c_{ll}}{D} a^i \right].$$

The update direction now has two components: one along l^i (the gradient direction) and one along a^i (the off-diagonal coupling direction). Depending on the choices of $\vec{a}$ and $\vec{b}$, this can lead to a variety of different optimization behaviors:

- If $\vec{a} = \vec{0}$ or $\vec{a} \propto \vec{l}$, we have plain gradient decent with an adaptive learning rate.
- If $\vec{a} = \text{momentum}$, we have a momentum-like update with an adaptive learning rate.
- If $\vec{a} = \vec{\theta}$ (the parameter vector), we have an update that combines gradient descent with a form of weight decay or growth depending on the sign of c_u/D .

For $\vec{a} = \vec{b} = \vec{l}$ the update simplifies to

$$\delta\theta^i = -\eta \frac{l^i}{1 + 3c_u}$$

which is similar to the diagonal case but with a stronger damping factor in the denominator ($\xi = 3$ instead of $\xi = 1$).

6 Implementation

All four families of optimisers (fixed diagonal, learnable scalar, learnable diagonal, off-diagonal) have been implemented in both JAX (via Optax) and PyTorch.

Files

Variant	JAX	PyTorch
Fixed diagonal	<code>jax_fixed.py</code>	<code>torch_fixed.py</code>
Learnable scalar	<code>jax_learnable_scalar.py</code>	<code>torch_learnable_scalar.py</code>
Learnable diagonal	<code>jax_learnable_diag.py</code>	<code>torch_learnable_diag.py</code>
Off-diagonal	<code>jax_offdiag.py</code>	<code>torch_offdiag.py</code>

General Update Direction

In all implementations, the goal is to realise the update

$$\delta\theta^i = -\eta g^{ij} r_j,$$

where r_j is the update direction. In ordinary gradient descent one would simply take $r_j = \partial_j \mathcal{L}$, but we can also test other choices such as momentum-based updates $r_j = m_j$ where m_j is an exponentially-smoothed average of past gradients (the usual “momentum” vector). This separation between the geometric direction l_i appearing in the metric and the update direction r_i allows the optimiser to use the metric to shape the descent direction in a more flexible way. Concretely, the four combinations we test are:

- $l = g$, $r = g$: pure geometric gradient descent.

- $\mathbf{l} = \mathbf{g}$, $\mathbf{r} = \mathbf{m}$: the metric is shaped by the instantaneous gradient but the step is taken along the momentum direction.
- $\mathbf{l} = \mathbf{m}$, $\mathbf{r} = \mathbf{g}$: the metric incorporates accumulated history but the step is raw gradient.
- $\mathbf{l} = \mathbf{m}$, $\mathbf{r} = \mathbf{m}$: both the metric and the step are momentum-based.

Off-Diagonal: Computational Cost

The off-diagonal induced metric is

$$G = \gamma \mathbb{1}_N + \xi(\mathbf{a}\mathbf{l}^\top + \mathbf{l}\mathbf{b}^\top + \mathbf{l}\mathbf{l}^\top),$$

where γ is a scalar (we take the base metric to be $\gamma \mathbb{1}_N$). We can write this as a rank-2 update,

$$G = \gamma \mathbb{1}_N + \xi UV^\top, \quad U = (\mathbf{a}, \mathbf{l}), \quad V = (\mathbf{l}, \mathbf{b} + \mathbf{l}),$$

so that by Sherman-Morrison-Woodbury,

$$G^{-1}\mathbf{r} = \frac{1}{\gamma}\mathbf{r} - \frac{\xi}{\gamma^2}U\left(\mathbb{1}_2 + \frac{\xi}{\gamma}V^\top U\right)^{-1}V^\top\mathbf{r}.$$

The contraction $G^{-1}\mathbf{r}$ costs $O(N)$: it requires only the inversion of a 2×2 matrix and inner products, never forming a full $N \times N$ matrix.

Off-Diagonal: Modes for $\mathbf{l}$, $\mathbf{a}$, and $\mathbf{b}$

The off-diagonal optimiser provides several built-in modes to choose the geometric vectors $\mathbf{l}$, $\mathbf{a}$, and $\mathbf{b}$.

The base direction $\mathbf{l}$: $\mathbf{l}$ encodes the direction along which the loss coordinate of the ambient space stretches. Two choices are provided:

- **Gradient mode:** $\mathbf{l}_i = \partial_i \mathcal{L}$. This reproduces the behaviour of the original diagonal construction, where the metric is shaped by the instantaneous slope of the loss.
- **Momentum mode:** $\mathbf{l}_i = \mathbf{m}_i$, the bias-corrected EMA of gradients. In this mode the metric is influenced not only by the current gradient but by the accumulated direction of travel over recent steps.

Either choice may be paired with $\mathbf{r} = \mathbf{g}$ or $\mathbf{r} = \mathbf{m}$ independently.

The off-diagonal vectors $\mathbf{a}$ and $\mathbf{b}$: Several built-in modes allow for quick experimentation:

- **Zero mode:** $\mathbf{a}_i = 0$ or $\mathbf{b}_i = 0$. This eliminates one branch of the off-diagonal coupling.
- **Gradient mode:** $\mathbf{a}_i$ or $\mathbf{b}_i$ is aligned with $\partial_i \mathcal{L}$, emphasising the local slope of the loss surface.
- **Momentum mode:** $\mathbf{a}_i$ or $\mathbf{b}_i$ is aligned with the momentum vector $\mathbf{m}_i$, incorporating information from recent history.
- **Parameter mode:** $\mathbf{a}_i = \theta^i$ or $\mathbf{b}_i = \theta^i$, coupling the geometry directly to the position in parameter space.
- **Same-as- $\mathbf{a}$ mode:** $\mathbf{b}_i = \mathbf{a}_i$, producing a symmetric off-diagonal structure.

For full generality, the implementations also accept user-provided vectors or functions that construct $\mathbf{a}$ and $\mathbf{b}$ dynamically at each optimisation step.

Learnable Variants: Metric State and Update

For the learnable scalar and diagonal variants, the metric parameters are part of the optimiser state and are updated at every step alongside the momentum buffer.

In the **scalar case** the state holds a single scalar s (initialised to zero, so $\alpha = e^s = 1$ at the start). At each step:

1. Compute $v = \xi e^s \|\mathbf{g}\|^2$ using the current gradients.
2. Update $s \leftarrow s + \mu v - \mu \lambda s$ (gradient ascent + L2 regularisation).
3. Clip $s \leftarrow \text{clip}(s, -s_{\max}, s_{\max})$.

The parameter update then uses the new $\alpha = e^s$ for the current step.

In the **diagonal case** the state holds a tensor $\mathbf{s}$ of the same shape as the parameters (initialised to zero). At each step:

1. Compute $v = \xi \sum_k e^{s_k} g_k^2$ (scalar denominator trace).
2. Update $s_i \leftarrow s_i + \mu \xi e^{s_i} g_i^2 - \mu \lambda s_i$ for each i .
3. Mean-centre: $s_i \leftarrow s_i - \bar{s}$.
4. Clip $\mathbf{s} \leftarrow \text{clip}(\mathbf{s}, -s_{\max}, s_{\max})$.

Both variants use the *same* gradient for both the metric update and the parameter update within a single step (there is no separate forward pass for the meta-gradient).

Computational Cost

All variants are $O(N)$ per step. No variant ever materialises a full $N \times N$ matrix:

- **Fixed diagonal**: element-wise multiply by the scalar γ .
- **Learnable scalar**: element-wise multiply by the scalar $\alpha = e^s$, plus an $O(1)$ scalar update for s .
- **Learnable diagonal**: element-wise multiply by e^s , plus an $O(N)$ element-wise update for s .
- **Off-diagonal**: inner products and a single 2×2 matrix inversion.

The metric learning overhead for the learnable variants is therefore negligible compared to the forward and backward passes of the model.

Shared Hyperparameters

The following hyperparameters are common to all variants:

- η (**lr**): base learning rate.
- β_m (**momentum**): EMA decay for the momentum buffer.
- ξ (**xi**): strength of the induced metric correction.
- β (**beta**): EMA decay for the denominator estimate $\hat{v}$.
- λ_θ (**weight_decay**): decoupled weight decay (AdamW style).

The learnable variants additionally require:

- μ (**metric_lr**): step size for the online metric update.
- λ (**metric_reg**): L2 regularisation on s or $\mathbf{s}$.
- $s_{\max}$ (**metric_clip**): log-domain clipping bound on s , constraining $\alpha = e^s \in [e^{-s_{\max}}, e^{s_{\max}}]$.

7 Results

All optimisers are evaluated together on the same suite of benchmark tasks. Results are compared head-to-head in the following notebooks:

- analysis/small_examples.ipynb — 2D synthetic functions (Beale, Rosenbrock, Himmelblau, Ackley, Rastrigin).




Figure 1: Small-example benchmark results for all optimiser variants. Each plot shows the best hyperparameter configuration found by the sweep. (*Results for the learnable scalar and diagonal variants to be added.*)